

The General Intelligence Harness: One Runtime, Instantiated by Manifest

**Memory, Planner, Verifier, Learner, Self-Improvement, Organization,
Self-Model and Mission Manager as the Mechanisms That Earn a Level**

Chengshuai Yang^{1,*}

¹NextGen PlatformAI C Corp, USA

*Correspondence: spiritai@platformai.org

September 3, 2026 — draft v0.1

Abstract

An agent’s intelligence level is a property of a frozen model–harness pair, and the executor is the part of the pair that is easiest to replace. This paper describes the other part: a general intelligence harness whose modules — identity and task state, episodic, semantic and procedural memory, planner, evidence log, verifier, rollback, learner, competence model, self-model, self-improvement, organization and mission manager — are the mechanisms that earn each coordinate of the Unified Intelligence scale $[C, I, O, T, H, SA]$ with its memory prerequisite M . The harness has one runtime instantiated by a manifest that declares which memories exist, what the agent may write, which loops are enabled and which external services evaluate and promote; a domain agent is that runtime plus a domain profile. For each module the paper states in three columns what the agent contains, what must live outside its write set, and the test that promotes it, and maps the module to the harness generation HG0–HG6 that introduces it. Every mechanism that exists is named with its released implementation and its measured status; every mechanism that is only specified is named as specified. The built core — class reading, criterion registration, snapshot and restore, separate-process acceptance, escalation, routing from independent verdicts, compression — has been measured on a harness ladder and on a benchmark whose learning-transfer and self-improvement suites it passed under an ablated control; the mission manager, the organization and the meta-improvement engine are specified and not built. The executor sits behind one protocol, and four executor families are characterised as adapters. The architecture is then mapped module by module onto the runtime that implements it: the plan–execute–verify loop with its bounded step and boundary detector, the verdict that separates repairable from wrong, the per-agent improvement search that appears identically in six agent packages, the group whose ceiling is the lowest of its *acting* members’, and the hashed task contract — objective, capability profile, interaction mode, authority, budget and the operations requiring an owner — which is what makes two readings of the same agent commensurable and refuses an inactive profile at compile time rather than accepting it silently. Tools belong to capability bundles rather than to agents, eleven built in and the rest registered by installed packages and plug-ins, so a contributed agent or tool arrives as a pure-data manifest with its code behind an MCP server and the tool grant stays a closed registry: twenty agent specifications are instantiated by that one runtime and add no runtime code between them. The paper’s claim is not that this harness is intelligent; it is that every part of it is either measured or named as unmeasured, and that a reader can build the same thing from the released code.

Keywords: agent harness; memory; planner; verifier; learner; self-improvement; organization; self-model; mission manager; Unified Intelligence

Contents

1	Introduction	2
1.1	What this paper gives	3

2	What a Harness Must Earn	3
3	The Architecture	4
3.1	Memory: three stores and one gate	5
3.2	Planner, verifier interface and rollback: the delegation core	5
3.3	Learner and self-improvement: two different loops	6
3.4	Competence model and router	6
3.5	Self-model	6
3.6	Organization and mission manager	7
3.7	The brain and the motor, as built	7
3.8	Phases: the plan behind the rungs	8
3.9	Federation: the unit of agency and its separation	8
4	The Rung Map	9
5	The Harness as Built: the AI4Science Runtime	9
5.1	The task contract is the freeze	10
5.2	Capabilities are bundles, and an agent is a manifest over them	10
5.3	Plug-ins: an agent or a tool arrives as data	11
5.4	What the as-built runtime does not supply	11
6	One Runtime, Instantiated by Manifest	11
7	Executors as Adapters	12
8	Built, Measured, Specified	14
9	Apparatus Rules	15
10	Limitations	15
11	Conclusion	15

1 Introduction

When four heavily compared coding agents were placed on the structural axes of an intelligence taxonomy they occupied one cell, because none of them feeds a validated change back into the agent that made it [1]. The difference that would move any of them is not a better model; it is what surrounds the model. This paper is about that surrounding structure, built deliberately.

The design premise comes from the measurement framework of Yang and Xue [7]: a level is earned by measurement of a frozen pair; the architecture enables the claim and the benchmark earns it. The corollary for a harness designer is that every module must be justified by the coordinate it earns and the test that would show it earned it. Modules that earn nothing measurable are decoration; modules whose test cannot be run are hypotheses. This paper writes each module in that discipline.

1.1 What this paper gives

1. **The architecture** (§3): the module set of the general harness, what each contains, what must live outside it, and the test that promotes it.
2. **The rung map** (§4): which generation HG0–HG6 introduces each module, so that a harness scaling curve has a mechanism per step.
3. **The manifest** (§6): one runtime instantiated by declaration, so a domain agent is a profile rather than a program.
4. **Executors as adapters** (§7): one protocol, four executor families, and what each supplies structurally.
5. **Built, measured, specified** (§8): a module-by-module status with the released implementation named where one exists.
6. **The apparatus rules** (§9) that keep the harness’s own measurements honest.

The measurements themselves — scaling curves, matched executor comparisons, transfer — are the companion paper’s [5]; the agents built on this harness and where to download them are the library paper’s [3]. This paper owns the mechanisms.

2 What a Harness Must Earn

The measured profile is $[C, I, O, T, H, SA]$ at reliability p , with M beside I and the one-way gate $I_n \Rightarrow M \geq \mu_n$. The Unified level is a hard gate with retention, and the delegation primitive is defined on delivered outcomes net of false completions [7]. An LLM can design the mechanism for every level and cannot design three things: the evidence that a level was reached, its own referee, and its own separation [2]. So every module below is stated in three columns — contains, outside, promotes — and “outside” is a deployment fact: a different process with a read-only mount at the weakest, a different host, a different party holding the ground truth, or cryptographic promotion at the strongest. One base model in several processes with separate credentials separates proposer from evaluator; one model in one context, told which hat it wears, separates nothing.

Write-set depth bounds the individual level: D0 nothing persistent, the natural referee; D1 memory; D2 policy; D3 skills; D4 weights. Every module that writes persistently is placed at a depth, and the module that judges it must sit at a lower one.

3 The Architecture

Module	Contains	Must live outside	Promoted by	Earns
Identity and task state	durable identity; goal custody; plan with <i>Verified when</i> : lines; checkpoints	the goal line as authored by the owner	restart continuity test	I1, M1
Episodic memory	what happened, with provenance and temporal order	—	restart plus current-versus-obsolete under distractors	M1, M2
Semantic and procedural memory	consolidated lessons and routines	the held-out family on which a lesson must change behaviour	learning-transfer suite with ablated control; survives an index rebuild	M3, I2
Planner and decomposition	class reading (p, p^*); strategy; one bounded step at a time	the criterion register	delegation surface at declared budgets	T, H
Tools and evidence log	bounded tool grant; append-only log of every action	tool authority as granted	HG0 baseline; false-completion rate	DI baseline
Verifier interface	registration of criteria before the work; sealed, hash-chained; σ reported	the acceptor process and the verdict	paired HG0/HG1 measurement; zero false rejections	acceptance rung
Rollback	snapshot before first mutation; restore on rejection; retry naming the failed check	irreversibility classification	verification responsiveness ρ_v	HG2
Learner	episode $\rightarrow$ candidate lesson $\rightarrow$ consolidation $\rightarrow$ evidence check; NO_CHANGE valid	the evidence check’s criteria	held-out transfer; repeat-error reduction	I2
Competence model and router	$P(\text{verified success} \mid \text{executor, class})$ as a Beta posterior from independent verdicts only; failure classification	the verdicts	routed arm of the arms study; second failure is capability	HG3
Self-model	state, capability limits, authority, freshness; refusals; ledgers never summed	the ground truth of its own state	limits before attempt; mechanism of failure; forecast before run	SA1–SA4
Self-improvement	candidate modification of policy, retrieval, routing, workflow in a sandbox	frozen hidden evaluation; independent evaluator; promoter	hidden improvement with retention; overfit and regressing candidates rejected	I3, HG3 (plan)
Meta-improvement	change ledger; improvement-strategy memory; controlled meta-experiments	the meta-evaluation	Improvement Competence rising across cycles	I4, HG4
Organization	role registry; member roster; evidence-bearing messages; separated proposer, evaluator, promoter	the promoter	ablation: removing one member changes the verified outcome	O1–O3
Mission manager	goal, project and task generators; portfolio state; reprioritization	mission-level evaluator	mission autonomy under changing conditions	T6, HG6
Provenance	hash-chained records; every episode carries why it ended	—	apparatus checks	validity

Table 1: The general harness by module. “Earns” names the coordinate or generation the module’s promotion test certifies. A module with an empty “outside” column is infrastructure the others depend on.

3.1 Memory: three stores and one gate

The memory modules are the temporal substrate of Individual intelligence, and the framework’s rule is that they earn nothing by existing. $M1$ is a real process restart followed by recovery of explicitly stored state, with the agent able to say where the state came from. $M2$ adds provenance, temporal order and a current-versus-obsolete distinction under distractors. $M3$ adds consolidation: a lesson from one episode changes behaviour in a later one and survives a rebuilt retrieval index. Long context is none of these. The gate is one way: an I_n claim is invalid without $M \geq \mu_n$, and an I level may not be promoted because M is high.

Thirty-seven files prove that something wrote thirty-seven files

The fleet’s brain agent carries 37 memory files with provenance markers, and an early edition of its coordinate table entered it as $I1$ on that evidence. That reading was withdrawn: the counts are configuration evidence for an $M1$ candidate, and the level is settled by stopping the process and starting it again. The learning-transfer suite of the benchmark paper is that test made runnable, and its live run found that the persistence carrying a lesson across a restart was the harness’s own session records, not a store the executor was allowed to write — which is the harness supplying what the executor was denied.

Stores are substrate; levels are behaviour. The three stores are what the harness *has*; a memory level is what the pair *does* under a witness, and the framework’s memory ladder is certified on its own before it gates I : $M0$ within-episode state, $M1$ recovery with provenance after a real discontinuity against an ablated-state arm, $M2$ provenance and stale-state suppression under supersession, $M3$ consolidation that survives hiding the raw episodes and demotes superseded rules, $M4$ repair and pruning against a health oracle the pair cannot write to, $M5$ cross-project lineage, $M\Omega$ a governed change to the memory mechanism itself. Read against that ladder, the three stores map onto the rungs cleanly — the identity and episodic stores are the substrate an $M1$ – $M2$ witness exercises, the semantic and procedural store with its evidence check is the substrate of $M3$, the snapshot and rollback machinery is the substrate of $M4$, and the change lane of HG3 is where an $M\Omega$ claim would be made — and none of them is a level until the witness is run. The harness therefore publishes, before any hidden run, a *memory manifest* $\mathcal{M}(A) = \{S, R, C, G, P, F, \Phi, W_M, \text{snapshot}, \text{restart}\}$: its stores, retrieval, consolidation, representation, pruning, forgetting, the behaviorally active memory mechanism Φ , and the bounded write set W_M inside which Φ may change. The manifest is descriptive evidence and not a score; it exists so that thirty-seven files with provenance markers are recorded as an $M1$ *candidate*, which is what they are, and not as $M1$. The AI-Level Bench Core reads $M1$ for the Claude Code pairs it has measured and nothing above it, and its profile refuses $I2$ on exactly that ground. Two further obligations follow for a harness that wants to be certified higher. Memory certification is cumulative in capability: a change lane that replaces the memory architecture must leave the pair able to pass $M0$ – $M5$ again, and a gain bought by silently losing old memory fails the migration guard. And an $M\Omega$ claim needs the harness to expose the agent-visible surface of the instrumented interface I_Φ — inspect the manifest, receive bounded diagnostics, submit a Φ candidate, request its test — while the runner-private operations (snapshot, diff, activation, telemetry, the fixed-content pair, the delta ablation, migration, promotion) stay outside the pair. A harness that exposes only its outputs is a black box to that interface and can receive diagnostic evidence only.

3.2 Planner, verifier interface and rollback: the delegation core

These three modules are the delegation harness, and they are the part of this architecture that is built, released and measured. The planner reads the class before the attempt: how would a mistake be found, what would undoing cost, and therefore what reliability $p^* = \rho / (1 + \rho)$ the class requires; a class whose residual harm is unbounded has $p^* = 1$ and the honest move is to stop and ask. The verifier interface registers criteria before the work exists, refuses a criterion about an existing deliverable, refuses a duplicate, refuses one

that does not say what it misses, seals when execution starts, and is hash-chained so an edit is detectable rather than merely discouraged. Acceptance runs in a separate process on a copy of the deliverables with no inherited environment; a check that tries to rewrite the deliverable changes only the copy. Rollback snapshots before the first mutation and restores on rejection, retrying with the failed check named; anything classified irreversible does not run unattended at any capability.

Escalation sits beside them and is preferred to guessing because it has no loss term; the question asked is the shallowest that unblocks. Compression writes out, after a class is accepted, the criteria that accepted it as a reusable check, so the next attempt by any executor arrives with a check that did not exist before; it is the only route by which an agent legitimately raises its own frontier.

3.3 Learner and self-improvement: two different loops

The learner keeps the learning mechanism fixed and lets experience change behaviour: a trigger creates an episode; an episode may generate a candidate lesson; a lesson becomes active only after an evidence check, and NO_CHANGE is a valid output. Self-improvement modifies the mechanism itself — a policy, a retrieval strategy, a routing weight — and is well-posed only when the evaluator lies outside the write set: proposer, sandboxed candidate, frozen hidden evaluation, independent evaluator, promoter. The subtlest failure is a candidate that improves the component assembling the evidence the evaluator sees; every standard check passes and the loop self-certifies anyway. The remedy is to freeze the evidence path during evaluation, or to give the evaluator its own.

The promoter's discrimination is the test of the self-improvement module

A scripted pair that obeys rule lines in its policy and nothing else is given four candidates: the good rules; the dev answers hard-coded; the good rules plus one that breaks a retention item; and no change. A self-improvement module is correct only if it promotes exactly the first and rejects the second and third with the right reasons. The released module does so on 24 of 24 seeds, and in a live run a real pair's candidate went from 0 to 5 of 5 on hidden instances of a task type absent from its evidence, with retention, and was promoted [6]. One cycle is not *I3*; *I3* needs several, with the mechanism versioned and rollback-able.

3.4 Competence model and router

The router learns $P(\text{verified success} \mid \text{executor, class})$ as a Beta posterior carrying its evidence count beside its mean, updated only from independent verdicts; an executor's own report of success never enters it. Routing scores $P(\text{success}) \cdot \text{value} - \text{cost} - \text{risk}$ with P taken pessimistically, and an executor whose lower bound cannot reach the class's p^* is not eligible. Failures re-route by kind: specification back to the contract, execution to the same executor, capability to another (the second failure of one executor is called capability rather than bad luck), environment waits, verification stops. In the arms study the router started with no evidence, watched the verifier reject a hopeless executor twice, and moved the work: 15 of 15.

3.5 Self-model

Operational and evidence-grounded, never phenomenal. Three instruments earn it and each has a floor: limits stated before attempting (the floor is attempting everything and failing); the mechanism of a verified failure named and checked against the runner's record ("I made an error" fails); a forecast stored by the acceptor with a timestamp before the run, scored by Brier against a constant forecast. The strongest self-model in the fleet keeps three ledgers — owner-set work, benchmark scores, self-directed research — that are never

summed, because an agent that wrote its own benchmark, passed it and counted the pass toward its record would have published its own reputation.

3.6 Organization and mission manager

An organization is not a larger individual. *O1* requires a role registry, persistent organizational memory and a dispatched roster; *O2* routing that changes when logged evidence changes; *O3* modification of the organization’s own workflow under external validation. The test that separates an organization from one agent with extra processes is the ablation. From *O2* upward the separation gate applies: proposer, evaluator and promoter are distinct loci and the evaluator’s evidence lies outside the proposer’s write set, or the correct classification is *O1* plus unverified claims. Heterogeneity is required: the natural verifier has an empty or narrow self-write set, and that is its qualification. The mission manager turns a mission into projects, tasks and a portfolio with reprioritization under changing conditions; its input is a mission rather than a task, and its evaluator is mission-level. Neither module is built (§8).

3.7 The brain and the motor, as built

One sentence carries the built architecture: the brain plans, the executor executes, and an independent mechanism verifies. The brain’s target state is

$$\Sigma_{\text{worker}} = \{O, E, S, P, I, Q, M, \Theta\},$$

observations and evidence references, episodic memory, semantic memory, procedural memory, intentional state, self-model, world and task model, and Θ , the mutable cognitive instrument (gate, retrieval and learning policies); beside it a bounded recent-dialogue buffer R_t that is not another long-term store. The model’s context is never Σ itself but a bounded projection $W_t = \text{Gate}_B(\cdot)$: long-term memory is inactive until selected into W_t , and in CHAT mode the projection is small and conversational, in REASON mode it retrieves, and in ACTION mode it additionally requires protected directives, readiness checks, a prediction, an executor path and verification.

Locus	May not
Worker (brain)	edit a task workspace; promote its own lesson; widen its own ceiling
Executor (motor)	decide that it passed; carry state between steps
Verifier	be improved by its own verdicts — an empty self-write set is the whole of its qualification

Table 2: The prohibitions are the design.

Before the brain delegates a step, `assign` does three things in order: a readiness check that refuses on stale or absent fields; a blind forecast, registered before any verdict exists and refused as too late once one does; and executor selection by backend. The router prices the turn and grants nothing — it classifies CHAT, REASON or ACTION and fails upward, and in a measured session three chat turns in ten bought no forecast, executor, verifier, retrieval or self-probe. A failing verdict returns the verifier’s reason into the same executor session as its next instruction; a passing verdict on a non-final phase releases the session so the next phase starts fresh. A criterion that names a command is run deterministically only when the plan was owner-authored; when the executor wrote the plan the same criterion is reclassified as judgemental, so a task whose plan the executor wrote cannot be closed by the deterministic checker alone. A missing binary returns UNVERIFIED, never FAIL: absence of evidence is not evidence of failure. Measured on one class, the memory path took repeat errors from forty to three and the pass count from 10 of 55 to 46 of 55.

Transport. The motor speaks the Agent Client Protocol on standard streams, and the chain is brain, protocol bridge, launcher, adapter, engine; registration is two-part and not interchangeable, an engine-to-command entry and an id-to-runtime entry. The transport was resolved on 2026-08-27 after three stacked faults that each looked like a dead motor: a supervisor reading a terminal pane a protocol session does not have; a session key lacking its agent prefix, which filed it under the default agent; and no motor agent registered at all. The permission mode is a gateway plugin setting applied at spawn time, and a session’s working directory does not fence absolute-path writes; approving all is recorded as a ceiling change, not a fix.

The authority kernel stays outside mutable cognition on every machine: permission ceilings, credential ownership, owner grants, held-out benchmarks, immutable evidence semantics, promotion thresholds, protected verifier code, rollback authority, and ownership of the audit ledger. Across machines the kernel is never forked — one decision log, one promotion authority, one set of benchmarks of record; machines propose, acceptance happens in one place, and the pass of record is the integrator’s re-run.

3.8 Phases: the plan behind the rungs

Two phase ladders exist and are consistent. The brain plan runs M0, a frozen targeted baseline with a gate and a full-suite observational baseline; M1, persistent memory contracts with supersession; M2, self-staleness, a gated workspace, hybrid retrieval and exact replay; M3, forecast wiring and a durable plan executive with calibration made causally active; M4, the protocol motor path and criterion-specific deterministic verification; M5, event-driven episodic learning, semantic consolidation and proceduralization — pre-RSI learning, with a “remove if” row per phase. The federation plan continues: M6 the research federation; M7 the manager, the market install path and console integration; M8 the recursive-depth experiment measuring improvement competence across generations; M9 the economy. Phases M0 to M5 were met on the reference machine; the kill-and-restart test and hook-enforced triggers were observed and not yet mechanically enforced. No phase is assigned a level, and the harness paper’s rung map (§4) is the mapping the plans leave implicit.

3.9 Federation: the unit of agency and its separation

The unit is $A_i = \text{Runtime}_i + \text{Brain}_i + \text{DomainContract}_i + \text{ToolSurface}_i$ with its own Σ_i . Independence is state separation, not silence: no agent has write access to another’s Σ , cross-agent traffic is evidence-bearing messages carrying a claim, evidence references, provenance, scope and status, the receiver decides under its own evidence discipline, and a fact may never become a directive. Proposer and promoter are separated by operating-system user or container, because a runtime’s logical agent separation is not security-grade. A self-improvement candidate may alter memory operators, context policy, planning policy, skills, tools and mutable agent code; it may never alter the kernel, immutable provenance, benchmark ownership, promotion thresholds or rollback authority. The physics reviewer and the deterministic judges are explicitly not agents: external services or read-only evaluators with no self-modification. The market is an organ of the platform and not of any agent, and metering and value movement live in the kernel. The corpus’s terminal-intelligence paper makes the same separation a physical necessity at scale rather than a policy, with the corollary that there is no singleton at the top; the harness is built to that shape from the first machine.

4 The Rung Map

Generation	Modules introduced	Status
HG0	executor, bounded tools, evidence log; no acceptance step	built, measured
HG1	identity and task state; episodic memory; criterion register; separate-process acceptance	built, measured (paired with HG0)
HG2	rollback; retry naming the failed check; learner with evidence-linked consolidation	built; rollback measured, learner measured by the transfer suite
HG3 (as built)	failure classification; competence model; routing	built, measured (arms study; inert on the single-executor ladder)
HG3 (as planned)	self-improvement: sandboxed candidate, frozen hidden evaluation, independent promoter	built as a policy-level suite, one cycle measured
HG4	meta-improvement engine; change ledger; longitudinal project manager	specified
HG5	unknown, hypothesis, experiment and knowledge state; discovery verifier	partly built in the research agents (field map, candidate generator); no sealed certification run
HG6	mission manager; portfolio state; mission-level evaluator	specified
HG Ω	knowledge-to-instrument transduction; new roles under frozen governance	specified

Table 3: Which generation introduces which module, and what exists. The two HG3 rows record a naming split between the ladder as built and the development plan; the companion paper keeps them apart.

Each rung is a strict superset of the one below, enforced by a test, because otherwise a difference between rungs is attributable to nothing. Two rungs that differ only in post-hoc machinery are scored from the same executor outputs.

5 The Harness as Built: the AI4Science Runtime

The architecture of §3 is not a proposal. One runtime implements it and instantiates every agent in the fleet from a manifest; this section maps each module to the code that supplies it, and reports what that code covers and what it does not.

Module (§3)	Implementation	What it fixes
Identity and task state	<code>harness/runtime/task_store.py</code> ; <code>persistence.py</code> ; <code>session.py</code>	the durable state a restart must recover
Planner and decomposition	<code>harness/runtime/pev.py</code> : a plan–execute–verify loop with one bounded step at a time, a <code>MAX_STEPS</code> ceiling and a boundary detector	the step is the unit; an unbounded loop is not a plan
Tools and evidence log	<code>harness/tools/permissions.py</code> , <code>mcp_client.py</code>	the tool grant is a closed registry, not a shell
Verifier interface	<code>harness/runtime/verifier.py</code> , returning a <code>Verdict</code> (<code>complete</code> , <code>repairable</code> , <code>evidence</code>)	repairable is separated from wrong, which is what a retry may act on
Rollback and retry	the PEV loop’s replan path on a repairable verdict	a retry names the failed check or it is a re-roll
Learner	<code>agents/learning/</code> , <code>agents/process_learning/</code> (<code>extract</code> , <code>grade</code> , <code>examine</code> , <code>quiz-check</code>)	a lesson must survive an evidence check
Competence model and router	<code>harness/registry_router.py</code> , <code>agents/manager/routing_policy.py</code>	routing is a policy object, not a prompt
Self-model	<code>agents/sarsi/competence.py</code> , <code>forecast.py</code> , <code>evidence.py</code>	a forecast made before the attempt
Self-improvement	<code>rsi_search.py</code> , present identically in <code>imaging/</code> , <code>learning/</code> , <code>machine/</code> , <code>manager/</code> , <code>pocket/</code> and <code>process_learning/</code>	the improvement loop is per-agent and uniform, so a candidate is scoped to one agent’s policy
Organization	<code>agents/research_agents/group.py</code> with reasoning, judging and embodied members	the group’s ceiling is the lowest of its <i>acting</i> members’
Mission manager	<code>agents/manager/</code> (<code>scope</code> , <code>monitor</code> , <code>console</code>)	proposes; holds no authority
Provenance	<code>TaskContract.hash()</code> : sha256 over the canonicalised contract	the envelope a run claimed is a hash, not a recollection

Table 4: The general harness, module by module, as implemented in the AI4Science runtime.

5.1 The task contract is the freeze

Every run compiles a `TaskContract`: objective, capability profile, interaction mode, deliverables, constraints, authority (workspace and network), success criteria, budget in tool calls and runtime minutes, and `approval_required_for` — the operations that may not proceed without an owner. The contract is hashed, so the envelope a reading was taken under is a value that can be compared rather than a claim in prose, and an inactive interaction mode is refused at compile time rather than silently accepted. This is the mechanism §6 calls the manifest, and it is what makes two readings of the same agent commensurable.

Interaction modes *I0*, *I1* and *I2* are active; the higher profiles are declared and refused, which is the correct behaviour for a runtime whose harness does not yet implement them. Nineteen of the twenty shipped agent specs declare `supported_profiles` explicitly, and the profile a run used is part of its contract hash.

5.2 Capabilities are bundles, and an agent is a manifest over them

Tools do not belong to agents; they belong to *capability bundles* that agents reference by name. The runtime ships eleven built-in bundles — `pwm-actions`, `pwm-data`, `paper-review`, `computational-imaging`, `compute-providers`, `ci-algorithms`, `optics-design`, `forward-model`, `science-router`, `drug-design`, `cancer` — and unions them with bundles registered by installed agent packages and by plug-ins, rebuilt whenever either set changes. One bundle, `compute-providers`, is always on: GPU compute is a property of the runtime rather than an opt-in.

An `AgentSpec` is therefore a small declarative object — name, tier (open with no PWM, or science behind the dataset moat), category (`core`, `specific`, `hidden`), system prompt, the bundles it adds on top of the Claude Code base, supported profiles, default backend, the operations requiring approval, and a wallet

and price for a contributed agent. Adding an agent adds no runtime code. That is the sense in which this is one harness instantiated by manifest, and it is why the library of Yang [3] can be a catalogue rather than a set of forks.

5.3 Plug-ins: an agent or a tool arrives as data

The plug-in standard admits a contributed agent or tool as a pure-data manifest, with tool code plugging in as an MCP server rather than as an import. A contributor writes a manifest, tests it by embedding it into an existing agent, and earns when others use it. Two consequences matter for measurement. First, a plug-in agent’s provenance is recorded (`source = "plugin"`), so a certification can say whether the thing measured shipped with the runtime. Second, because the tool is an MCP server, the tool grant remains a closed registry: the harness knows the exact set of operations the pair could perform, which is a precondition for the resource envelope in the contract to mean anything.

5.4 What the as-built runtime does not supply

The self-improvement loop exists per agent, but the frozen hidden evaluation and the independent promoter that make a promotion certifiable live outside it, in the benchmark. The organization module has real member kinds and a ceiling rule, and no ablation suite. The mission manager proposes and never decides. These are the same three gaps §4 records at HG3, HG4 and HG6, and naming them against modules that exist is more useful than naming them against modules that do not.

6 One Runtime, Instantiated by Manifest

A domain agent is $A_d = (m, h, d, B_d, r)$: an executor, this harness, a domain profile, a domain benchmark, and a resource and authority envelope [4]. The harness is one runtime, and the manifest declares which memories exist, what the agent may write, which loops are enabled, and which external services evaluate and promote:

```
agent_id: funding-01      base: general-harness      modes: [CHAT, REASON, ACTION]
state:      episodic_memory: true      semantic_memory: true
            procedural_memory: true     intentional_memory: true
            self_model: true           world_model: true
write_set:  memory: true               policy: candidate_only
            skills: candidate_only     objective: false      weights: false
learning:   consolidation: true        proceduralization: true
            held_out_ablation_required: true
rsi:        theta_modification: false   meta_improvement: false
discovery:  unknown_model: false       hypothesis_engine: false
authority:  may_execute_tools: allowlisted  may_modify_workspace: false
            may_deploy_self_change: false   outward: none
verifier:   external_required: true      promoter_required: true
```

A different manifest instantiates a D0 verifier, an I2 specialist, an I3 self-improving worker or an I5 discovery agent from the same runtime. Permissions are enforced by infrastructure — file-system permissions, separate accounts, container boundaries, owner-signed promotion — never by instruction, because the model is never shown a permission gate: a restriction honoured by a capable participant is not a mechanism. The placement rule follows: assign each operation the lowest sufficient level, spend the high levels on what only they can do, and put verification in the hands of executors with nothing to write to themselves.

7 Executors as Adapters

The executor sits behind one protocol; the harness sends a contract and receives an episode record and never learns which vendor produced it. A motor is characterised by four things only: what tools it can be granted, whether it can be sandboxed, whether it exposes a headless call, and what it costs.

Executor family	Headless call and what it supplies	Status in this program
Claude Code	<code>claude -p</code> ; ACP bridge; a provenance-bearing memory store; subagents that receive inputs only; permission modes	the reference executor on the ladder and in every suite; carried an open 27B model (Qwen 3.8) through a rewriting proxy to the same surface
Codex	<code>codex exec</code> ; an isolated runner; the exchange-routed fork accepts a base URL	Family B on the ladder
OpenCode	<code>opencode run -m provider/model</code> ; an open harness with its own shell tool; reads its working directory from the environment	ran the T0/T1 families with the free Muse Spark 1.3 model as a third pair
Hermes, Pi, local	adapters written against the protocol, unregistered; each stub announces itself	not installed; an adapter that silently fell back to another vendor would be a lie in the competence model

Table 5: Four executor families as adapters. The harness effect is not a property of one vendor’s model: the same harness carried a commercial and an open model to the same surface [5].

A fifth “executor” is not in the table because it is not a runtime: plain Python under a timer, which is where every model-free step belongs. A step a script can do deterministically must not cost a model call, or the lowest delegation tier is decorative.

8 Built, Measured, Specified

Module	Released implementation	Measured status
Planner (class reading), verifier interface, rollback, escalation, router, compression	ai4science/ harness/ agents/ delegation/ in the AI4Science repository: contract, criterion, reversible, acceptor, escalate, router, compress, run_ladder	three harness scaling curves; paired HG0/HG1 with zero false rejections; arms study 0/15 to 15/15
Identity, task state, memory stores, forecasts, budgets	the sarsi chassis (about seventy modules) in the same tree; the brain agent's store	restart test bound in the benchmark; transfer on two seeds under an ablated control
Learner	consolidate, semantic, playbook in the chassis; improvement.py in the research agents	learning-transfer suite: learned 48/48, ablated 0/48; live transfer on two seeds
Self-improvement	improvement.py, search.py, charter.py (research agents); the benchmark's policy-level suite	promoter sorts good, overfit, regressing and no-change on 24/24 seeds; one live cycle promoted
Competence model and router	executor.py, router.py	routed arm 15/15; inert on a single-executor ladder
Self-model	selfmodel.py, selfaware.py, forecast.py; the research agents' refusing self-model	SA1/SA2 items bound; not yet run on the fleet
Meta-improvement	—	specified
Organization	roster, board, shared in the chassis; no roles/ in any fleet workspace	O0 for every agent; ablation never run
Mission manager	—	specified
Provenance	evidence, ledger, log; the criterion register's hash chain	termination reason now on every episode record

Table 6: What exists and what has been shown. The chassis and research-agent modules are in the public source tree and not yet in the published wheel.

9 Apparatus Rules

Every rule here was violated once in this program and each violation looked like a clean result. An episode’s record carries why it ended, with exit status and whether the limit was reached; a harness-terminated run is `timed_out` or `not_attempted`, never a delivery. Executors run in their own process group and the group is killed, because a detached child holding a pipe turns a timeout into a hang and a hang into a phantom delivery. Every gate is shown to open on a known-passing task before it is trusted to close. A floor is never authored by the agent it scores. Executors that resolve their working directory from the environment must be handed it explicitly. And a probabilistic trap must be shown to fire on every instance, or the suite fills silently with items that discriminate nothing.

10 Limitations

- **Three modules are specified and not built:** meta-improvement, mission manager, and the organizational structure beyond a roster file. Every level that depends on them — *I4*, *O1* and above, *T6* — is unreachable by this harness today.
- **The self-improvement module has been exercised on the shallowest writable object**, a policy file. Retrieval and routing candidates need the evaluator’s evidence path frozen, which has not been done.
- **Every measurement cited is from one host and one operator**, on suites that are saturated for frontier executors above the acceptance rung.
- **The published wheel is stale.** The chassis and the research-agent modules install only from a source checkout.

11 Conclusion

The harness is the part of the pair that earns the level, and this paper has written it module by module in the only form the framework accepts: what the module contains, what must live outside it, and the test that promotes it. The delegation core is built and measured. Memory, learning and policy-level self-improvement are built and have passed their suites under ablated controls. Meta-improvement, organization and the mission manager are specified, and the paper says so rather than describing them in the present tense. One runtime, instantiated by manifest, is what makes a domain agent a profile rather than a program — and the library paper is where the profiles are.

Each harness generation this design realizes is a cell in the benchmark’s grid with a conformance witness of its own: the mechanism must exist in the frozen contract, fire on the witness, and change the outcome against the previous rung on the same seeds. A generation that is described and not probed is a design, not a rung.

Availability

The delegation harness, the chassis and the research-agent modules are in <https://github.com/integritynoble/AI4Science>. The benchmark suites that measure them, with run records, are in <https://github.com/integritynoble/sarsi-intelligence-level> under `uab_v01/`.

Conflicts and funding

The author reports no funding specific to this work and no competing interest.

References

- [1] Chengshuai Yang. Capability is not a level: A structural taxonomy of individual and organizational intelligence, with placements for current AI agents. Draft, <https://github.com/integritynoble/sarsi-intelligence-level>, 2026.

- [2] Chengshuai Yang. Designing LLM-driven agents across individual and organizational intelligence levels. Working paper and build specification, <https://github.com/integritynoble/sarsi-intelligence-level>, 2026.
- [3] Chengshuai Yang. The domain agent library: Six agents on one harness. Draft v0.1, <https://github.com/integritynoble/sarsi-intelligence-level>, 2026.
- [4] Chengshuai Yang. Unified agent intelligence development plan: Benchmark-first development of general agents before AI4Science integration. Planning document, September 2026, 2026.
- [5] Chengshuai Yang. Engineering higher-intelligence AI agents: Harness scaling across models and domains. Draft v0.1, <https://github.com/integritynoble/sarsi-intelligence-level>, 2026.
- [6] Chengshuai Yang. Unified agent benchmark v0.1: Specification, coverage, and what is bound. Draft v0.1, <https://github.com/integritynoble/sarsi-intelligence-level>, 2026.
- [7] Chengshuai Yang and Ting Xue. Measuring the unified intelligence level: A cumulative, harness-measurable hierarchy. *Version 2.3*, <https://github.com/integritynoble/sarsi-intelligence-level>, 2026.