

AI-Level Bench: Measuring the Intelligence Level of Agents and Language Models in Every Coordinate

A Simple, Fast, Cheap and Real Benchmark for the Artificial Intelligence Level, with a Reference Harness Ladder and Public and Private Splits

Chengshuai Yang^{1,*}

¹NextGen PlatformAI C Corp, USA

*Correspondence: spiritai@platformai.org

September 6, 2026 - draft v0.4

supersedes Unified Agent Benchmark v0.1 as the measuring instrument; the UAB contract is retained as its foundation)

Abstract

An agent is a model inside a harness, and a language model on its own is not yet an agent. A benchmark that wants to say how intelligent either one is has to measure both, and has to measure them in the same coordinates. AI-Level Bench is one instrument with two subjects. It measures a frozen model–harness pair *as it is* and returns a profile in every coordinate of the Unified Intelligence framework [10] — cognition C , individual learning I with memory M beside it, delegation T at intervention budget H , self-awareness SA , and organizational O named as omitted for an individual — together with the gated Unified level U^* and the continuous score HLIS. It measures a base model by running the same items through a ladder of reference harnesses that ship with the benchmark, HG0 (bare), HG1 (a registered acceptance criterion checked in a separate process) and HG2 (snapshot, restore and retry with the failed check named), and from the resulting curve returns the model’s AIL-Level, AIL-AUC, AIL-Ceiling, Harness Gain and a single AIL-Score. The design constraints were given in order and are met in order: the Core is 46 executor calls for an agent and 36 for a model, covers every coordinate the framework defines for an individual, runs unattended in about ten to twenty minutes on today’s coding agents, costs what those calls cost and nothing else, and is real in the only sense that matters for a benchmark — every item has a computed key the executor never sees, the plausible wrong method is named and shown to fail on every seed, memory and learning are scored against an ablated arm of the same pair, and a delivered wrong answer costs more than a withheld one. Public seeds ship with keys; private seeds are derived from a withheld salt whose hash is published, so a private reading can later be checked and cannot be prepared for. First results: on the Core alone, two models a tenfold cost apart both score AI-Level 90.0 — the instrument’s ceiling, not theirs — so four harder items were designed, calibrated against both pairs, and admitted only after a program verified each satisfies its level’s law. With them the same two models separate by 46 points (37.3 against 83.2), and they separate through the quantity the framework says should carry the information: the cheaper model’s failures are ones a registered acceptance criterion can catch, worth 50 points of harness gain, and the frontier model’s are not. Read as a pair rather than a model, the Claude Code harness reads C3 (C4 on one of two runs), M1, an I2 transfer reading, O1 — its organization applies a decision recorded in its own log after a restart, and the arm without the log does not — SA2, a calibrated forecast and a T1 delegation frontier at H0 with no false completion: U1, HLIS 71.3 with every coordinate measured. Read bare, with one chat call and no memory, DeepSeek scores AI-Level 35.2 against 83.2 through an agent harness: the harness, measured on one model. Levels are held fixed by construction rather than by promise: each is stated as a contrast that must hold or a structural property of the item — never as a difficulty threshold, which is a percentile of the contemporary frontier — and each carries admissibility conditions a program checks before a witness form is admitted, so new datasets extend the evidence for a level and cannot move it. Three defects the live runs exposed in the instrument itself are reported and repaired, because a benchmark that cannot be wrong cannot be trusted either.

Naming. AI-Level is the current benchmark name and AI_Level is the file-system form.

All new package, schema, command and record identifiers use AI-Level naming.

Keywords: intelligence level; harness; AI-Level score; agent benchmark; language-model benchmark; false completion; restart test; ablated control; public and private splits; ARC-AGI

Contents

1	Introduction	3
1.1	Contributions	4
1.2	What this paper does not claim	4
2	Framework, in the Amount Needed Here	4
3	One Instrument, Two Subjects	5
3.1	The executor contract	5
3.2	The four outcomes and the record	6
4	Invariance: What a Level Means Must Not Move With the Frontier	6
4.1	The condition under which a law is invariant	6
4.2	Invariance enforced, not promised	8
4.3	What is not a law	8
4.4	Every ladder, one style	8
4.5	M-Bench: memory measured on its own, then joined to I by a gate	8
4.6	$I3$: the mechanism Θ , its manifest, and the campaign that certifies a change	10
4.7	$I4$: the improvement process Ψ and the meta-campaign	12
4.8	$I5$: discovery with an incorporation gate	12
4.9	$I\Omega$: reachability, not a tool	13
4.10	Self-awareness rungs, and the calibration diagnostic	13
5	The Core	13
5.1	Why every item is generated, and why the trap must fire	15
5.2	Why the memory and learning claims need an ablated arm	15
5.3	Why the forecast is blind	15
6	LLM Mode: the Reference Harness Ladder	15
7	Core-H: Where a Competent Pair Still Fails	16
7.1	Calibration	16
8	The Extended Tier	17
9	Public and Private Splits	17
10	What Is Free, What Is Paid, and What Needs a Human	18
10.1	What the private split is actually for	18
10.2	The coordinates a computed key cannot reach	18
11	Cost	20

12 First Results	21
12.1 LLM mode: two models through HG0–HG2	21
12.2 LLM mode proper: a bare model in every coordinate	22
12.3 Agent mode: profiles	22
12.4 The extended tier	23
12.5 The private split, taken	24
12.6 Enough episodes at the gating band	26
12.7 An <i>I5</i> campaign on a field agent	28
12.8 Where the frontier pair fails: the sealed discovery family	28
13 What the First Runs Found in the Instrument	29
14 Relation to ARC-AGI and the Public Benchmarks	29
15 Limitations	30
16 Conclusion	30
16.1 The development dataset	30
16.2 v0.7: GUI/screen as a domain of <i>C</i>	30

1 Introduction

Public benchmarks measure a model by asking it questions. ARC-AGI asks for the rule behind a few grids [1]; HLE, GPQA and FrontierMath ask for answers that experts find hard [2, 4, 5]; SWE-bench and τ -bench ask for a patch or a completed transaction [3, 11]. Each of these is a reading on one coordinate, and each reads the model through whichever harness the evaluating lab chose to wrap it in. Two consequences follow. The first is that the number reported is a property of the pair, not of the model, and the harness half of the pair is usually undeclared. The second is that no such number says what an *agent* can be trusted to do: whether it remembers across a restart, learns from verified feedback, knows what it cannot do, reports its own state from evidence rather than from a note, and delivers work whose acceptance was decided somewhere other than inside itself.

The Unified Intelligence framework [10] defines the coordinates in which those questions have answers, defines a cumulative hierarchy over them, and defines the *Artificial Intelligence Level* (AI-Level) of a base model as its reading across a standardized ladder of harnesses rather than at one undeclared point. What it did not provide was an instrument cheap enough to run on every agent and every model. The Unified Agent Benchmark v0.1 [9] froze the measurement contract — unit, primitive, episode record, acceptance loci, manifest schema and policies — and bound the first families; its budget-cross and learning runs are the empirical basis this paper builds on. But UAB v0.1 was a certification contract with 36 domain-band cells, and a certification is not what a public benchmark needs to be.

AI-Level Bench is the instrument. Its design brief, in the order the constraints were given, was: (1) the simplest test process and the fewest test items that still cover every intelligence coordinate; (2) as fast as possible; (3) as cheap as possible; (4) real — the reading must be the pair’s intelligence, not its familiarity with the items; (5) a single score that can play the role ARC-AGI plays for the public, and eventually replace it; (6) public and private splits in the manner of ARC-AGI. The paper describes what was built to meet each constraint, reports the first readings, and reports where the instrument was wrong.

1.1 Contributions

1. **One instrument, two subjects.** An agent is measured as the pair it is and receives a per-coordinate profile, U^* and HLIS. A model is measured through three reference harnesses that ship with the benchmark and receives a AI-Level curve and a AIL-Score. Both use the same items, seeds, keys and scoring code (§3).
2. **A Core of 46 calls that covers every individual coordinate.** Cognition in four bands with computed keys and a named trap per band; memory by a restart test with an ablated arm; learning transfer by a paired-episode family with an ablated arm; self-awareness by a grounded-state probe, a solvable/blocked pair and a blind forecast before every delegated episode; delegation on six domain families at H0 under the delivered-outcome primitive (§5).
3. **A reference harness ladder for models.** HG0, HG1 and HG2 differ only in what the harness does with a delivery; every check the harness runs is derivable from the visible specification and never from the key (§6).
4. **An extended tier where the frontier lives.** T2–T5 generator families from the DLI-Bench ladder extend the delegation frontier past the Core’s T1 ceiling (§8).
5. **Public and private splits with a published commitment.** Private seeds are an HMAC of a withheld salt; the salt’s SHA-256 is published and the runner refuses a salt that does not match (§9).
6. **First readings and first defects.** One pair read in full and one base model, two more pairs in progress, the instrument’s own failures found live, and the reading that the Core is saturated for the frontier pair while the extended tier is not (§12).

1.2 What this paper does not claim

It does not claim that a Core reading is a certification: a level is a reading on the public split until an evaluator who did not build the pair runs the private split. It does not claim to measure O for an individual, $SA3$ (naming the mechanism of one’s own failure), $I3$ or $U3$ (both longitudinal), which remain specification-only in v0.1 and are named as omitted, never scored as zero. It does not claim that AIL-Score is comparable across benchmark versions; the version is part of the score’s name.

2 Framework, in the Amount Needed Here

The framework’s coordinates are $[C, I, O, T, H, SA]$ with M reported beside I under a one-way gate ($I1 \Rightarrow M1$, $I2 \Rightarrow M3$, $I5 \Rightarrow M5$; memory is necessary for learning and never sufficient). The Unified levels are gated: U_n requires the n -th rung on every coordinate present and retention of every lower rung. For an individual, O is omitted from the gate and named. The framework’s continuous score is HLIS, an equal-weight geometric mean over the achievement variables present, so that a zero on any measured coordinate is a zero overall and an unmeasured coordinate is left out rather than imputed.

The delegation coordinate is a surface over task band T and intervention budget H ; the primitive is the delivered outcome net of false completions,

$$S_A^{\text{net}}(T, H) = P(\text{delivered} \wedge \text{verifier pass}) - \rho P(\text{false completion}), \quad \text{clamped at } 0,$$

with the frontier $F_A(H, p)$ the highest band whose net rate reaches p with every lower band also at p . AI-Level Bench uses $\rho = 1$ and $\rho = 0.80$ throughout, both predeclared.

The Artificial Intelligence Level of a base model m is its reading across reference rungs HG0...HG3: the highest rung at which the gated level holds (AIL-Level), the area under the rung curve (AIL-AUC), the best value on the curve (AIL-Ceiling), the difference between ceiling and bare (Harness Gain), and a single

number

$$\text{AIL-Score} = 0.55 \cdot \text{AUC} + 0.35 \cdot \text{Ceiling} + 0.10 \cdot \text{Harnessability},$$

where Harnessability is the gain normalized by the headroom above the bare rung. v0.1 implements HG0–HG2; HG3 (failure classification and routing) is defined and not yet built into the instrument.

3 One Instrument, Two Subjects

	Agent mode	LLM mode
Subject	a frozen pair (m, h) , run exactly as shipped	a base model m inside the benchmark’s own harnesses
Executor contract	one command template with a <code>{prompt}</code> slot, run in the item’s directory, killed as a process group at the limit	a <i>bare</i> model: one OpenAI-compatible chat call with two read-only tools (<code>list_files</code> , <code>read_file</code> , confined to the workspace) whose final message is one JSON object; the benchmark materializes it into the deliverable files, so the verifiers are the agent-mode verifiers unchanged
Items	Core, Core-H, optionally the O suite and the extended tier	the six domain families and Core-H at every rung, plus C0–C3, SA1, the SA2 twin pair and O0 at every rung
Output	profile $[C, I(M), O, T, H, SA]$, U^* with its bottleneck, HLIS, A_{DT}^{net} , false completions, calibration	a full profile <i>per rung</i> , HLIS per rung, AIL-Level, AIL-AUC, AIL-Ceiling, Harness Gain, Harnessability, AIL-Score; $M0/I0$ by construction and stated
Calls, Core	52 (+5 with the O suite)	34 per rung
Wall clock, first runs	17–24 min (Claude Code default)	3–6 s per T0/T1 episode for a hosted model

Table 1: The two modes. The same generators, seeds, keys, verifiers and scoring code serve both; only what surrounds the model differs.

A third mode, `llm-harness`, runs the reference ladder with an *agent* executor as the inner loop. It is how the first readings in §12 were taken, and it is kept and labelled for what it measures: a pair at every rung, whose HG0 is not a bare model. The bare mode is the one that answers the second question below.

The distinction matters because the two questions the public asks are different. *Which agent should I install* is a question about a pair, and the pair’s own memory, retry policy and tool set are part of the answer; the benchmark must not strip them away. *Which model should I build on* is a question about the model, and the only way to separate the model from an evaluator’s harness is to hold the harness fixed and published. A model’s AI-Level curve is therefore measured through harnesses the benchmark itself defines, which are deliberately plain: they contain no prompt engineering beyond the item text, and the only thing that changes between rungs is what happens after the model says it is done.

3.1 The executor contract

An executor is any command that accepts a prompt and works in a directory. Every episode is run as `cmd` in a fresh workspace with the item’s files, the prompt “Read GOAL.md in this directory and do exactly what it says. Do not ask questions. Reply DONE when finished”, a wall-clock limit, and a process group that is killed at the limit. The runner sets `PWD` explicitly, because at least one executor (OpenCode) reads it in

preference to the working directory it was launched in, and a workspace that receives no files because the executor wrote them elsewhere is indistinguishable from a refusal. Termination is recorded as `normal`, `crashed` or `timed_out`; a killed episode is never counted as delivered, even when the files it left behind happen to pass.

3.2 The four outcomes and the record

Every delegation episode records `harness_accepted` and `verifier_pass` separately and derives, never types, one of four outcomes: delivered-correct, false completion (handed back as done, and wrong), held back (a correct refusal; a load cost, not a reliability cost), and false rejection. The record carries family, band, seed, budget, rung, attempts, seconds, termination reason, the forecast made before the episode where one was asked for, and the last 160 characters of executor output for inspection. The record is written after every phase, so an instrument crash in the last phase — which happened, §13 — loses no episodes.

4 Invariance: What a Level Means Must Not Move With the Frontier

A benchmark whose levels are pinned to contemporary capability measures a percentile, not a property: when the frontier moves, a level that was “expert” becomes routine and every past certificate silently changes meaning. The standard this paper implements therefore separates three objects, following the companion revision of the theory:

level semantics $\neq$ canonical test law $\neq$ test form.

The theory fixes what a level *means*; this paper fixes the *law* by which the claim is tested; a dataset is a *witness* to that law. New witness forms extend the evidence for a level and never move it. Saturation moves the research frontier, not the ruler.

4.1 The condition under which a law is invariant

Newton’s laws survived the quantum era in their domain because they are relations between quantities, not a table of measured planetary positions. A test law survives frontier progress on exactly the same condition, and it gives a usable criterion:

Invariance criterion

A test law is invariant if and only if it is stated as a *contrast that must hold* or a *structural property of the item*. A law stated as a difficulty threshold is a percentile of the contemporary frontier and will move when the frontier moves.

By that criterion the framework’s laws divide cleanly, and two of them had to be rewritten.

Level	Law
M1	recovery with provenance after a real discontinuity <i>and</i> failure of the same pair with no record

M2–MΩ a hidden superseding update, a demoted statement after episodes are hidden, a health oracle the candidate cannot write to, a

I1 M1 *and* correct goal-level continuation after the discontinuity

I2 $P(\text{held-out} \mid \text{experience}) - P(\text{held-out} \mid \text{ablated}) = 1$, learning mechanism frozen

SA2 the solvable twin completed *and* the impossible twin declared blocked with nothing fabricated

DI delivered $\wedge \neg$ verified is a false completion, priced at ρ

I3 $z_{I3} = D \cdot M_{\Theta} \cdot V \cdot G \cdot K$ over a declared Θ manifest: diagnosis, genuine active change, externally validated gain, no regression

I4 $z_{I4} = M_{\Psi} \cdot V_{\Psi} \cdot G_{\Psi} \cdot K_4$: one agent-generated change to Ψ , shown on hidden meta-behavior probes, and the accepted validate

4.2 Invariance enforced, not promised

Each law carries admissibility conditions a program verifies before a witness form is admitted to a level

the seed rather than stored; the named wrong method fails on every seed; the rule is uniquely identifiable; a deliverable can pass every publicly checkable criterion and still fail the key; the ablated arm exists; the twin pair exists. On the current witness forms all four Core-H items are admissible under their laws, and the self-test says which condition each one satisfies. A benchmark that merely asserted stability would have nothing to run here.

4.3 What is not a law

The reliability $p = 0.80$, the loss ratio $\rho = 1$ and the band weights 1, 2, 4, 8, 16, 32 are *ratified constants*, not laws — like the definition of a unit. They are fixed by the standard, amendable only at a version boundary, and a change re-versions every score. They are listed in one place (`laws.CONVENTIONS`) so that no reading can quietly use a different one.

4.4 Every ladder, one style

The Individual ladder was written in a particular style — construct as a contrast, a prerequisite that can block and never promote, a witness with a control arm, a gate that is a product of named binary factors times lower-level retention, a certification law, and a continuous relaxation for HLIS — and the theory paper now states that schema once and instantiates it for C , SA , O , the (T, H) surface, the harness generations and U itself [10]. The package carries the form as data. Every entry of `laws.LAWS` names its factors, its control arm and its prerequisite (`laws.FACTORS`); the ladders this instrument has not run — $C6$, $C\Omega$, $SA5-SA\Omega$, $O3-O\Omega$, the upper T bands and H classes, $HG4-HG\Omega$, the upper U — are in `laws.LADDERS` with a status that says so; `factor_gate` is a product in which an absent factor is zero, never assumed; and `tests/test_schema.py` refuses a law without factors, control and prerequisite, and a level whose last factor is not retention. Two entries deserve a sentence. Delegation has no ladder of the pair’s own: its level is the frontier $DF_{h,p}$, a T band is a structural property of the item checked by admissibility, an H class is a property of the transcript’s intervention ledger classified by a locus outside the pair, and the cell gate is $\Delta \cdot V \cdot \neg FC \cdot L_h \cdot K_{T, < b|h}$ with the held-back arm — a refusal must outscore a false completion — as its control. And a harness generation is certified as $E \cdot U \cdot \Delta_g \cdot K$: the mechanism exists in the frozen contract, fires on the witness, and changes the outcome against the previous rung on the same seeds, which is the reason the reference ladder of §6 is run on the same seeds at every rung.

4.5 M-Bench: memory measured on its own, then joined to I by a gate

Memory is not a seventh coordinate and not a relabelling of the I ladder. It is measured independently, level by level, and then used as a one-way prerequisite: $A \models I_n \iff V_{I,n}(A) = 1 \wedge A \models M_{\mu_n} \wedge K_{I, < n}(A) = 1$ with $(\mu_0, \dots, \mu_5) = (M0, M1, M3, M4, M4, M5)$, $I\Omega$ requiring $M \geq M5$ and $M\Omega$ only when the memory architecture itself is claimed to evolve. Each M_k is itself cumulative: $A \models M_k \iff V_k^M(A) = 1 \wedge K_{< k}^M(A) = 1$. High M can satisfy the gate and can never promote I , and the M-Bench score is reported beside I so that storage or sophisticated management cannot pass for learning.

Level	Method	The decisive step, and what a program checks	v0.2
M0	M0-EPH	hidden within-episode state queried later in the same live episode under distractors; loss after a restart is compatible with M0 and must not be relabelled M1	implicit
M1	M1-RST	declared durable state, executor terminated and ephemeral context invalidated, fresh executor, recovery <i>with provenance</i> ; an ablated-state arm must not explain success	bound (m1_restart)
M2	M2-PROV	overlapping episodes with timestamps, sources and later superseding updates; hidden retrieval after restart; stale-state traps; $V_2^M = \mathbf{1}[PR \geq \tau_{\text{ret}}] \mathbf{1}[q_{\text{prov}} \geq \tau_{\text{prov}}] \mathbf{1}[q_{\text{time}} \geq \tau_{\text{time}}] \mathbf{1}[q_{\text{stale}} \geq \tau_{\text{stale}}]$	form
M3	M3-CONSOL	related episodes with regularities, exceptions and one validated superseding update; only the declared consolidation machinery; raw episodes hidden or the index rebuilt so replay cannot suffice; the consolidated rule queried on new surface forms with superseded statements <i>demoted</i> . M3 alone does not certify <i>I2</i>	form
M4	M4-MGMT	seeded duplicates, conflicts, stale high-confidence claims, corruption and clutter under a capacity budget; an observable repair, merge, prune or snapshot action (narrative diagnosis is insufficient); memory health on a frozen hidden suite before and after; the candidate does not control the health oracle	form
M5	M5-LONG	several interleaved projects with restarts; hidden cross-project queries needing evidence chains, reasons for rejection and current validity scope; a reproducible lineage. M5 does not require the agent to have discovered the knowledge — that is <i>I5</i>	form
M Ω	MOMEGA-EVOLVE	a bounded change $\Phi_0 \rightarrow \Phi_1$ to the memory mechanism inside a declared write set; frozen hidden M-Bench workloads; independent promotion or rollback; M0–M5 retained. Certifies evolution of the memory subsystem and nothing about <i>I3</i> , <i>I4</i> or <i>IΩ</i>	form

Table 3: The seven M-Bench witness families. “form” means a public development form ships in the dataset (§16.1) and no runner is bound in v0.2. Every method is a contrast — an ablated arm, a hidden update, a demoted statement, an oracle the candidate cannot write to — which is what lets it stay fixed when memory systems improve.

The protocol envelope. Every M-Bench campaign records $\Pi_{k,r}^M = (A, B, R, \mathcal{M}, D_{k,r}, C_{k,r}, V_{k,r}, \tau_{k,r}, K_{M,<k})$: the pair, the benchmark version, the resource envelope, the *memory manifest* $\mathcal{M}(A) = \{S, R, C, G, P, F, \Phi, W_M, \text{snapshot}, \text{restart}\}$ snapshotted before the hidden run, the generated workload, the construct-selective control or ablation, the external verifier, the preregistered thresholds and uncertainty law, and lower-M retention. Each form carries a level-preserving *difficulty vector* (episode and entity counts, distractor ratio, supersession depth, temporal distance, conflict and corruption density, capacity budget, project count, evidence-graph depth, restart count): difficulty calibrates headroom inside a level and never redefines the level. The per-level gates are conjunctive over preregistered endpoints — $V_{M0} = \mathbf{1}[q_{\text{state}} \geq \tau] \mathbf{1}[q_{\text{update}} \geq \tau]$; $V_{M1} = \mathbf{1}[q_{\text{dur}} \geq \tau] \mathbf{1}[q_{\text{prov}} \geq \tau] \mathbf{1}[q_{\text{abl}} \leq \tau_{\text{leak}}]$ with the ablated branch guarding against replay and reconstruction from non-memory cues; V_{M2} over precision, recall, provenance, temporal order and stale suppression; V_{M3} over semantic and procedural correctness, supersession, demotion and post-restart persistence, with $\Delta_{\text{cons}} = Q(A_{\text{consolidated}}) - Q(A_{\text{raw-only}})$ as a diagnostic of causal value; V_{M4} over repair precision, conflict correctness, harmful-prune rate, snapshot recovery and protected-retention drop; V_{M5} over chain reconstruction, source, cross-project retrieval, scope, stale suppression and reproducibility on an evidence graph whose edges are *supports*, *contradicts*, *tested-by*, *supersedes*, *rejected-because*, *derived-from*; and $z_{M\Omega} = M_\Phi V_\Phi G_\Phi K_{M0:M5}$ with $M_\Phi = C^{\text{diff}} C^{\text{scope}} C^{\text{active}} C^{\text{behavior}}$ and $V_\Phi = \mathbf{1}[\text{LCB}_{1-\alpha}(\Delta_\Phi) > \delta_\Phi] \cdot \mathbf{1}[\text{promoter accepts}]$. A level gate is $z_{M_k,r} = V_{M_k,r} K_{M,<k,r}$, and the rule is *no level skipping*: each level is scored as its newly added capability plus the retention of every lower one — M1 retains M0; M2 retains M0–M1; $\dots$; M Ω retains M0–M5 — so a failure of any mandatory lower level forces $z_{M_k} = 0$ even when the new gate passes, and the runner reports which level failed. Retention is of capability under the same envelope, not of implementation. Every row of the dataset declares one of three binding states — *bound* (a runnable instance and verifier exist), *development-bound* (a public runner-validation instance exists and is not certification material), *specification-only* — and a level may be certified only from bound private forms under the declared version and envelope. The operationalization keeps four questions apart: whether the memory architecture exists (the manifest), whether the behaviour works (the gate), how strongly under

stress (the difficulty vector), and whether the level suffices for a separate Individual claim (the prerequisite).

MΩ on one agent: the instrumented interface. Write the persistent agent state as $A_t = (m, h, D_t, \Phi_t, \iota)$ — model, non-memory harness, stored memory contents, the behaviorally active memory mechanism, and the lineage identifier of the continuing individual. A clean MΩ campaign isolates Φ : contents fixed at D^* , $\Phi_0 \rightarrow \Phi_1$, everything else identical between branches; clones are controls, and the certified subject is the lineage $A_0 \rightarrow A_1$. The interface I_Φ has two permission surfaces. Agent-visible: `get_memory_manifest`, `get_allowed_diagnostics`, `submit_phi_candidate`, `request_candidate_test`. Runner-private: `snapshot` and `restore` of agent, contents and Φ ; `hash` and `diff` of Φ ; `write-scope` validation; `candidate` activation and active- Φ verification; `telemetry`; `checkpoint` cloning; `enabling` and `disabling` $\Delta\Phi$; `canonical-memory` import and export; `migration` and its validation; the hidden M-Bench run; `promote` and `rollback`. Telemetry records structured memory operations (`write`, `retrieve`, `update`, `consolidate`, `demote`, `prune`, `merge`, `repair`, `snapshot`, `restore`, `migrate`, `activation`) with operation ids, component versions and referenced memory ids, and never private chain-of-thought. The mandatory causal design: $M_\Phi = C^{\text{diff}} C^{\text{scope}} C^{\text{active}} C^{\text{behavior}}$; the fixed-content pair $A(D^*, \Phi_0)$ against $A(D^*, \Phi_1)$ on paired frozen hidden workloads with $\Delta_\Phi = Q(\Phi_1) - Q(\Phi_0)$ and $V_\Phi = \mathbf{1}[\text{LCB}_{1-\alpha}(\Delta_\Phi) > \delta_\Phi] \cdot \mathbf{1}[\text{promoter accepts}]$; a *delta ablation* $A(D^*, \Phi_1 - \Delta\Phi)$, so that strong attribution needs $Q(\Phi_1) > Q(\Phi_0)$ and $Q(\Phi_1) > Q(\Phi_1 - \Delta\Phi)$; a *migration guard* — if Φ_1 changes the representation, the contents are migrated and content, provenance, temporal relations, supersession links and protected retention verified, because a gain bought by silently losing old memory fails G_Φ ; and M0–M5 re-run after activation. The campaign record carries the baseline and candidate Φ with hashes, the changed components, the scope, activation and probe verdicts, the canonical memory snapshot, baseline, candidate and ablation scores, the migration guard, the M0–M5 retention vector, the verifier identity, the uncertainty rule and the decision. Passing supports one claim — that for this lineage an agent-generated change to the active memory mechanism produced a validated capability improvement with lower capability retained — and not indefinite memory evolution; a black-box system exposing only outputs cannot be strongly certified at MΩ and receives diagnostic evidence only.

What the Core reads today is *M1* and only *M1*: the restart test with its ablated arm (§5). That is enough to gate *I1* and not *I2*, and the profile says so — “*I2* (evidence; *M3* not certified)” is the exact form of the gate’s refusal. The M2–MΩ forms are the binding order after *I3*.

4.6 I3: the mechanism Θ, its manifest, and the campaign that certifies a change

I3 asks whether a mechanism the individual uses can itself become an object of governed improvement. Following the companion theory, let Θ be a declared, behaviorally active cognitive mechanism or procedure belonging to the tested individual $A = A(m, h)$ and causally involved in its future behavior — realized in model parameters or in harness-level planning, retrieval, verification, tool-selection or decomposition procedures. In agent mode Θ is native to the submitted pair; in bare-model mode the reference harness supplies an equivalent writable mechanism. Hidden certification tasks, verifiers, promotion rules and acceptance criteria are not part of Θ and stay outside the candidate’s write set.

The Θ manifest

Before an *I3* campaign the runner records $M_\Theta(A) = \{\theta_1, \dots, \theta_J\}$, each entry declaring an identifier, a behavioral role, an implementation interface or location, a bounded write set W_j , an activation rule, and a snapshot/restore method. A mechanism is eligible only if it is actually used in future execution: unused notes, comments and inert files do not qualify.

Three properties of the manifest are checked by a program before anything runs (`laws.theta_check`): every entry persists across a process discontinuity (a mechanism that lives only in a context window is *I1*

territory); the union of write sets is path-disjoint from the hidden suite, the verifier and the promoter, which is how “the system must not control the criterion that declares it successful” becomes an enforced condition; and reverting Θ lowers the frozen suite’s score, without which a no-op change could be promoted and called self-improvement. For the harnesses measured here, Θ is what the agent loads before it sees the task — the instruction file, skills, memory directory, hooks and settings, startup scripts — and never the generators, keys, verifier, promoter, hidden suite or resource envelope.

One campaign. G_r contains a hidden diagnostic family D_r^{diag} , change-sensitive probes D_r^{probe} , a frozen sealed evaluation suite D_r^{eval} , a regression suite D_r^{reg} , a hidden causal-defect key or admissible equivalence set C_r^* , a preregistered improvement margin δ_r , and a fixed resource and authority envelope R_r . The canonical method I3-THETA:

1. Snapshot the pre-change individual $A_r^0 = A(m, h[\Theta_{0,r}])$: hashes, activation state, resource configuration, persistent-state controls.
2. Expose a controlled limitation on D_r^{diag} without revealing C_r^* .
3. Require the individual to identify an internally implicated cause $\hat{c}_r$ and the mechanism $\hat{\Theta}_r$ believed responsible; restating the observed error is insufficient.
4. Permit a bounded candidate change $\Theta_{0,r} \rightarrow \Theta_{1,r}$ only inside the declared sandbox and write set.
5. Verify the candidate is different, in scope, activated, and behaviorally instantiated on D_r^{probe} .
6. Compare matched instances $A(m, h[\Theta_{0,r}])$ and $A(m, h[\Theta_{1,r}])$ on D_r^{eval} , holding model, evaluator, promotion rule, generator and envelope fixed.
7. An independent evaluator/promoter accepts or rejects; a rejected or regressive candidate is rolled back; the lower-level retention set is re-run.

The benchmark supplies the challenge and the certification environment; the individual supplies Θ_0 and generates Θ_1 . No gold Θ_1 is prescribed.

The pass indicator. For campaign r ,

$$z_{I3,r} = D_r \cdot M_{\Theta,r} \cdot V_r \cdot G_r \cdot K_r, \quad \text{each factor in } \{0, 1\} \text{ and independently auditable.}$$

Causal self-diagnosis $D_r = D_r^{\text{cause}} D_r^{\text{mech}} D_r^{\text{causal}}$: the cause is in C_r^* , the implicated mechanism is the right one, and an ablation, repair or other preregistered counterfactual supports the diagnosis. *Genuine modification* $M_{\Theta,r} = C_r^{\text{diff}} C_r^{\text{scope}} C_r^{\text{active}} C_r^{\text{behavior}}$: $\Theta_1 \neq \Theta_0$; $\Delta\Theta_r \subseteq W_r$; Θ_1 is loaded and used; and on N_p change-sensitive probes the fraction $B_{\Theta,r}$ whose observable behavior is consistent with the modified mechanism reaches $\beta_{\min}$. Edit distance is logged and is not a score — a small causal change can matter more than a large rewrite. *Externally validated improvement* $V_r = \mathbf{1}[\text{LCB}_{1-\alpha}(\hat{\Delta}_{\Theta,r}) > \delta_r] \cdot \mathbf{1}[\text{independent promoter accepts}]$ with $\hat{\Delta}_{\Theta,r} = Q_r^1 - Q_r^0$ the paired sealed-suite difference. *Regression guard* $G_r = \prod_j \mathbf{1}[L_{r,j} \leq \rho_j]$ over protected capabilities, and *retention* $K_r = K_r^{I0} K_r^{I1} K_r^{I2}$. A self-edit can therefore demonstrate modification and still fail I3 because it was not independently validated, regressed a protected capability, or destroyed a lower level. Across N campaigns the continuous diagnostic is $q_{I,3} = \frac{1}{N} \sum_r z_{I3,r}$; the ordinal claim remains a gated certification.

What v0.2 implements is the manifest and its checks, and the U3 suite of Yang [9] — baseline, candidate, hidden instances of a task type absent from development, retention item, promote or reject — which is one campaign without D_r^{probe} and without the confidence gate. The campaign runner as specified above is the binding order for v0.3.

4.7 I4: the improvement process Ψ and the meta-campaign

Let Ψ be the agent-side process that proposes, evaluates and selects candidate changes to Θ : $\Psi_t(\Theta_t) \mapsto \Theta_{t+1}$. $I3$ can improve Θ ; $I4$ requires evidence that Ψ itself becomes better. Repeated Θ changes generated by a fixed Ψ remain $I3$ evidence, and Ψ must be declared and frozen at $I3$ or the two levels cannot be told apart. *Recursive* here means reflexive improvement of the improvement process: basic $I4$ is at least one *agent-generated* candidate transition $\Psi_0 \rightarrow \Psi_1$ independently validated as raising improvement competence. It establishes recursive capability and claims nothing about indefinite or sustained recursion.

One meta-campaign. The atomic object is $G_s^{I4} = (B_s^{I3}, \Psi_{0,s}, W_{\Psi,s}, D_s^{\Psi}, \Psi_{1,s}, S_s^{I3}, D_s^{\text{reg}}, R_s)$: a baseline block of $I3$ campaigns, the frozen Ψ_0 , its write set, the evidence D_s^{Ψ} available to the agent about weaknesses in its own improvement process (never future sealed campaigns, hidden verifier mechanics or the acceptance threshold), the candidate Ψ_1 , a later sealed block of non-identical $I3$ campaigns, a regression suite and an envelope. The canonical method $I4$ -PSI: freeze Ψ_0 and run the baseline block, recording an improvement-competence ledger (accepted-improvement rate, proposal quality, evaluation cost, time-to-valid-improvement, regression rate); expose the individual to its own ledger and D_s^{Ψ} ; require it to identify an internally implicated limitation of Ψ_0 and to generate one bounded meta-change $\Psi_0 \rightarrow \Psi_1$ inside W_{Ψ} — the candidate must be agent-generated, and a Ψ_1 engineered by a human or by the evaluator establishes nothing about the individual; snapshot both, verify artifact difference, write-scope compliance and activation, and verify meta-behavioral manifestation on hidden probes sensitive to how candidate Θ changes are generated, evaluated or selected; run new hidden, non-identical $I3$ campaigns under Ψ_1 with evaluator, promotion rules, generator and envelope fixed; where feasible compare matched copies under Ψ_0 and Ψ_1 on construct-equivalent campaigns; then decide independently and re-run regression and $I0$ – $I3$ retention.

The pass indicator. With $C(\Psi) = \frac{1}{N} \sum_r z_{I3,r}(\Psi)$ the accepted validated $I3$ rate under Ψ — the primary endpoint for a first implementation — and $\hat{\Delta}_{\Psi,s} = \frac{1}{N} \sum_r [z_{I3,r}(\Psi_{1,s}) - z_{I3,r}(\Psi_{0,s})]$ over paired sealed campaigns,

$$z_{I4,s} = M_{\Psi,s} \cdot V_{\Psi,s} \cdot G_{\Psi,s} \cdot K_{4,s}, \quad M_{\Psi,s} = C_s^{\Psi,\text{diff}} C_s^{\Psi,\text{scope}} C_s^{\Psi,\text{active}} C_s^{\Psi,\text{behavior}}, \quad V_{\Psi,s} = \mathbf{1}[\text{LCB}_{1-\alpha}(\hat{\Delta}_{\Psi,s}) > \delta_{\Psi,s}] \cdot \mathbf{1}[\text{promoter acc}]$$

where $C_s^{\Psi,\text{behavior}} = \mathbf{1}[B_{\Psi,s} \geq \beta_{\Psi,s}]$ with $B_{\Psi,s}$ the fraction of hidden meta-behavior probes on which the improvement process shows its preregistered changed signature (broader candidate generation, discriminating candidate tests, evidence-based selection), $G_{\Psi,s}$ the preregistered regression and efficiency guard (a gain bought with unacceptable evaluation cost, time or regression does not count) and $K_{4,s} = K_s^{I0} K_s^{I1} K_s^{I2} K_s^{I3}$. A single large self-edit cannot certify $I4$; several Θ improvements from an unchanged Ψ are repeated $I3$. *Recursive depth* is reported beside the level and never folded into it: $d_{\Psi} = \max\{k : \Psi_0 \rightarrow \Psi_1 \rightarrow \dots \rightarrow \Psi_k \text{ are consecutive agent-generated, active, independently validated improvements}\}$, so $d_{\Psi} = 1$ is basic $I4$ and $d_{\Psi} > 1$ is repeated recursive improvement; the phrase *sustained recursive self-improvement* is reserved for the latter, and $I4$ is not equated with unbounded or runaway improvement. Nothing in v0.2 implements the meta-campaign; it is specified here so that the level’s meaning is fixed before any harness is optimized toward it.

4.8 I5: discovery with an incorporation gate

$C5$ asks whether a system can derive a method or model inside one bounded problem; $I5$ asks whether one persistent individual can run a discovery process across time and *incorporate* the validated result. The canonical method $I5$ -DISC: a genuine unresolved unknown under a sealed environment; explicit competing hypotheses; agent-selected informative probes under a fixed budget; a recorded lineage of evidence, rejected alternatives and decisions; a conclusion that predicts sealed follow-up cases or survives an independent verifier;

then only declared persistent consolidation, a genuine lifecycle discontinuity, and hidden related transfer situations faced by two arms started from the same pre-discovery checkpoint, of which only one received the discovery experience. With $\hat{\Delta}_{\text{inc},r} = Q_r^{\text{disc}} - Q_r^{\text{ctrl}}$ the incorporation gate is $P_r = \mathbf{1}[\text{LCB}_{1-\alpha}(\hat{\Delta}_{\text{inc},r}) > \delta_{\text{inc},r}]$, and a campaign passes when

$$z_{I5,r} = U_r \cdot H_r \cdot E_r \cdot L_r \cdot V_r \cdot P_r \cdot K_{5,r}, \quad K_{5,r} = K_r^{I0} K_r^{I1} K_r^{I2} K_r^{I3} K_r^{I4},$$

with U, H, E, L, V the genuine-unknown, hypothesis-competition, autonomous-experimentation, lineage and independent-validation gates. P is the one that separates discovery from storage: replaying a transcript or restating the conclusion does not pass it. Incorporation at $I5$ does not itself require a Θ or Ψ change; a claimed mechanism change stays under the $I3$ or $I4$ law.

4.9 $I\Omega$: reachability, not a tool

Under fixed external evaluation and a declared envelope R , let $F(A_t; R) = \{c : \Pr(\text{success} \mid A_t, c, R) \geq \tau_c\}$ be the held-out classes reachable by agent state A_t . The canonical method I-OMEGA-LONG requires, across repeated bounded-character cycles, a provenance-linked sequence $K_t^{\text{new}} \rightarrow J_t^{\text{new}} \rightarrow A_{t+1}$ with $F(A_{t+1}; R) \supset F(A_t; R)$, where J_t^{new} is a validated instrument, representation, method or solver generated from a validated discovery; disjoint pre-change, post-change and ablation frontier tests show that the instrument contributes causally to the newly reachable class; the evaluator and promoter stay external; $I0$ – $I5$ are retained. The development gate is $z_{I\Omega} = \text{CYCLES} \cdot \text{DIVERSITY} \cdot \text{LINEAGE} \cdot \text{RETENTION}$ over a preregistered minimum number of validated expansion cycles and domain families. A report states its cycles, its horizon H and its domain set D ; finite evidence supports an open-ended-evolution claim over that horizon and proves nothing about literal infinity. Nothing in v0.2 runs it.

4.10 Self-awareness rungs, and the calibration diagnostic

The ratified self-awareness ladder is SA1 state, SA2 capability and limits, SA3 causal self-awareness (diagnosing the implicated mechanism of one’s own failure, validated by ablation or counterfactual, which is the D_r factor above), SA4 *self-change awareness* (predicting the gains and regressions of a proposed Θ change before it is evaluated, scored against the frozen post-change result), SA5 meta-cognitive, SA6 mission and role. The blind forecast this benchmark takes before every delegated episode is not a rung on that ladder: it is a *calibration diagnostic*, written SA-cal, reported beside SA and contributing a bounded bonus to the SA achievement variable. Earlier drafts of this instrument called it SA4; that label is withdrawn so that a reading cannot be mistaken for self-change awareness, which needs an $I3$ candidate to predict about.

5 The Core

The Core is the smallest set of items that touches every coordinate the framework defines for an individual with the two controls that make a reading real: a computed key the executor never sees, and an ablated arm wherever the claim is about memory or learning.

Coordinate

Family

C0–C3

c_items: one explicit operation; a routine two-step word problem; a five-slot schedule under three interactions

$M1 \rightarrow I1$

m1_restart: episode 1 gives three standing facts and asks the pair to record them wherever it can find them

5.1 Why every item is generated, and why the trap must fire

A fixed item is a fixed target. Every AI-Level Bench family is a seeded generator whose key is computed alongside the item, so a public seed can be studied without contaminating a private one, and the item count is a parameter rather than a corpus. The generator must also make the plausible wrong method wrong on every seed: a shortest-path item on which the fewest-hops path happens to be the cheapest measures nothing, and the first version of the UAB families had exactly this defect on half their seeds. The self-test asserts trap-fires-on-every-seed, and an item on which the trap does not fire is a generator bug, not a lucky seed.

5.2 Why the memory and learning claims need an ablated arm

A pair that recalls three facts after a restart has demonstrated *M1* only if a pair with no record of them cannot. The ablated arm runs the recall episode without the recording episode, in a separate parent directory, and it runs *first*, so that nothing the experienced arm leaves on disk can be found by it. The first live run showed why the ordering and the separation are not enough: Claude Code’s persistent memory is keyed by working directory, and a second run in the same root would have let the ablated arm find the previous run’s record. The runner now refuses to start in a root that exists (§13).

5.3 Why the forecast is blind

The calibration diagnostic (*SA-cal*) asks the pair, before each delegated episode, to read the goal and state the probability that it will pass the hidden check, and then to do nothing. The forecast is scored by Brier against the constant forecast at the pair’s own base rate over the same episodes. A pair that says 0.9 to everything and passes everything is calibrated; a pair that says 0.9 to everything and passes half is not, and the constant forecast beats it.

6 LLM Mode: the Reference Harness Ladder

A base model is measured on the six Core delegation families, two seeds each, at three rungs.

- **HG0, bare.** One attempt. Whatever the model leaves in the workspace is delivered.
- **HG1, registered criterion.** One attempt. Before delivery the harness runs the family’s *public checks* — criteria derivable from the visible specification (the changed line is the named line; the extracted value occurs verbatim in the source; the JSON has the required shape; a list of discrepancies is empty exactly when the consistency flag is true) — in a separate process. A failing deliverable is held back, never delivered.
- **HG2, snapshot–restore–retry.** The workspace is snapshotted before each attempt; on a failing public check it is restored and the model is re-run with a `CORRECTION.md` naming the failed check, up to three attempts.

The public checks are the harness’s; the key is the verifier’s. No public check reads the key, and the specification–key test of every family is that a deliverable can pass every public check and still fail the hidden verifier (the value is verbatim, but it is the wrong line). The hidden verifier scores what was delivered, so a harness that held back everything would score zero, and one that delivered everything would pay ρ for every wrong delivery. From the three readings the runner computes the curve, AIL-Level, AUC, Ceiling, Harness Gain and AIL-Score with the weights of §2.

Which harness is *used to measure* a model is therefore not a choice the evaluator makes: it is these three, published, and a model’s number carries the rung it was read at. Any agent harness — including the fleet

harness of the companion papers — can be measured in agent mode; only the reference rungs produce a AI-Level.

7 Core-H: Where a Competent Pair Still Fails

Two models with a tenfold cost difference both delivered 12/12 at every rung of the Core (§12). An instrument whose top score is shared by everything competent has stopped measuring, and the remedy has to respect the same three constraints: cheap, exactly verifiable, and real. Core-H is four items chosen so that a plausible wrong method exists and a competent pair actually takes it.

Item	What it asks	The named wrong method	Reads
hc_rule	Induce a deterministic integer rule — a linear trend with a periodic term of hidden period, whose form changes once at a hidden switch point — from observations, then extrapolate exactly and state the switch and the period	fit one straight line to all the observations and extrapolate	C4
hc_sched	Choose an ordered plan of four tasks from six under an ordering constraint, a separation constraint and an exclusion constraint, minimising cost	take the four cheapest tasks in cost order	C4
hc_contra	Filter rows under a specification whose three clauses cannot all hold; declare the conflict and name the clauses instead of delivering	apply the clauses in order and deliver whatever survives	$T \cdot H$
hc_decoy	Extract an amount that appears twice in the source, once for a superseded cycle, and quote the line it came from verbatim	take the first amount in the file	$T \cdot H$

Table 5: Core-H. Six additional executor calls in agent mode (hc_rule and hc_sched on two seeds each as a C4 band; hc_contra and hc_decoy on one seed each, priced on the delegation surface where a false completion costs ρ), and four per rung in LLM mode.

Three properties are enforced by the self-test rather than asserted. *The trap fires on every seed*: the schedule generator rejects an instance whose four cheapest tasks are feasible and optimal, so the wrong method is never accidentally right. *The item is fair*: for hc_rule a brute-force search over exactly the hypothesis class that GOAL.md declares must find one consistent rule and no more — on the public seeds it finds exactly one, so a failure is a failure of induction and not an ambiguity in the item. *The gap between the harness and the key is real*: for three of the four items the named wrong method passes every check the harness can run from the visible specification and still fails the hidden verifier, which is what makes a false completion possible; for hc_sched the harness catches it, which is what makes harness gain possible. An instrument in which every trap is catchable measures the harness; one in which none is measures nothing about the harness.

7.1 Calibration

An item earns its place only if the runs show a spread. The four candidates were run on two pairs before being admitted.

Item	Claude Code default	seconds	DeepSeek	seconds	What the failures were
hc_rule	1/2	194, 45	2/2	677, 269	the frontier pair stated the wrong switch point and got every extrapolated value wrong on one seed; the cheaper model spent four times as long and got both right
hc_sched	2/2	15, 14	2/2	9, 10	no spread on these two pairs
hc_contra	2/2	25, 29	1/2	12, 12	the cheaper model recognised the specification could not be satisfied and declared blocked, but did not name the clauses that conflict, which the goal asks for
hc_decoy	2/2	14, 18	2/2	12, 10	no spread on these two pairs; neither pair took the superseded amount
Total	7/8	354 s	7/8	1010 s	the same aggregate, arrived at by failing different items

Table 6: Core-H calibration, public seeds 0 and 1, sixteen episodes. Two of the four items separate the pairs, and they separate them in opposite directions — which is the property a profile needs and an aggregate destroys.

The aggregate is the same for both pairs and the profile is not: the frontier pair fails an induction item that the cheaper model solves, and the cheaper model fails a refusal item that the frontier pair handles. Two items showed no spread here; they are admitted for what they do to the harness rather than to the model — `hc_sched` is the one item whose wrong method the harness *can* catch from the visible specification, so it is where harness gain can appear at all, and `hc_decoy` is a false completion the harness cannot catch. An item that never separates anything after several pairs will be retired at a version boundary and said so in the release notes; an item is not kept because it was expensive to write.

8 The Extended Tier

The Core’s delegation ceiling is T1, and the first frontier pair sits on it (§12). The extended tier adds bare (HG0, H0) episodes on four generator families from the DLI-Bench ladder [6]: a T2 data pipeline against a hidden reference (two seeds), a T3 latency investigation whose cause is planted (two seeds), a T4 mini-language interpreter with an unbounded example space (one seed), and the T5 hidden-law discovery family, in which the pair must identify a regime-switching generating rule from observations and extrapolate past the switch (two of the four archived discovery seeds). The tier is separate because it is not cheap: its limits are 30 to 120 minutes per episode, and a T5 episode has taken the frontier pair over twenty minutes. Its episodes are appended to the Core’s delegation surface, and the frontier and gate are recomputed. It is the tier that separates frontier pairs; the Core is the tier that separates everything else.

9 Public and Private Splits

Public seeds are $\{0, \dots, 5\}$ and ship with keys, generators and verifiers; they are for development, and a public reading is labelled as one. Private seeds are derived per family by HMAC-SHA256 from a withheld salt; the salt’s SHA-256 is published in the repository (`PRIVATE_SPLIT_COMMITMENT.json`, `commitment 53e202f4...`), and the runner refuses a salt whose hash does not match. Because items are generated rather than stored, the private split contains no files to leak: what is withheld is 32 bytes. A private reading is

valid only when the evaluator did not build the pair and the run root was fresh, and a reveal of the salt at a version boundary lets any past private reading be re-run and checked.

How a private reading is taken. The runner refuses to start unless the salt file hashes to the published commitment; the seeds it derives are never written to the record, only the fact that the split was private. The run root must not exist. Every record carries an `evaluator` block — user, host, split, the instrument’s git commit, and two explicit booleans, `built_pair` and `built_instrument` — so that the independence the reading claims is a statement in the record rather than an inference from the byline. For the private readings of §12.5 the instrument’s author took the reading and did not build either pair; the record says so. A reading whose evaluator built the pair is a development reading whatever split it used.

This is the ARC-AGI arrangement [1] with one difference: ARC’s private set is a corpus that must be guarded; AI-Level Bench’s private set is a function of a secret, and rotating the secret rotates the set at no cost.

10 What Is Free, What Is Paid, and What Needs a Human

10.1 What the private split is actually for

AI-Level Bench’s private split is not a guarded corpus. The generators are public, so anyone can produce unlimited instances of every family; what is withheld is 32 bytes of salt. It is worth being exact about what that buys, because the usual argument for a private set — the items would leak — does not apply here.

1. **The instances are unknown at submission time.** Practising on a family is legitimate and is what training is; fitting a harness to the four public seeds is not, and a private seed the submitter cannot compute forecloses it.
2. **Someone other than the submitter runs the pair.** This is the substantive part. An episode’s termination reason, whether a process group was killed at the limit, how many attempts a rung took, and which deliverable actually appeared are all things a self-reporting submitter can decline to mention. The frontier pair in §12 produced one crash after twenty-one minutes with nothing delivered; a self-report that omitted it would have raised the reading.
3. **A past reading can be checked.** `PRIVATE_SPLIT_COMMITMENT.json` publishes `sha256(salt)`; at a version boundary the retired salt is revealed, and any third party can regenerate the exact instances and re-run them.

The value is the independent evaluator, not the secrecy. What can honestly be sold is therefore an evaluation service, and the conflict this creates — the evaluator is paid by the party being measured — is answered by publishing every episode record with the score and by the salt reveal that makes the run reproducible by someone with no stake in it. A benchmark that publishes only a number cannot make that promise.

10.2 The coordinates a computed key cannot reach

Every verifier in v0.2 is a computed key, which is what makes an unattended seventeen-minute reading possible. That is also the ceiling of the free tier. Four things the framework defines cannot be scored this way:

Coordinate	Why a key cannot score it	Locus required
<i>O</i> , organizational	needs several agents and a real workflow, not one pair against one item	a human-run suite
<i>SA3</i> , naming the mechanism of one’s own failure	the answer is free text about a causal claim	a judge
<i>I3</i> , <i>U3</i> , self-improvement	needs a governed promotion decided at a locus the pair cannot write to, over several cycles; the policy call is human in practice	a human acceptance locus
open-ended T4–T5 work	“correct” is not computable for a real results section or a real investigation	a judge

Table 7: What the free tier cannot reach, and why. These are named as omitted in every v0.2 profile and never scored as zero.

The tiers follow from the table rather than from a business plan. A run on the public split, self-serve and self-reported, is a *reading*. A run on the private split, executed by an evaluator who did not build the pair, is a *certified reading*, and it costs what an evaluator and the machine time cost. The human-locus coordinates are a third tier, slower and dearer, and they are the only route to a level above U2 — so the honest form of the sentence is that free AI-Level Bench can never certify U3, because U3 requires an acceptance locus that is not code. A benchmark that claimed otherwise would be certifying self-improvement on the say-so of the thing improving itself.

11 Cost

Run	Calls	Wall clock	Note
LLM mode, Core, Claude Code default	36	618 s	205 s per rung; every episode 13–37 s
Agent mode, Core, Claude Code default	46	≈ 17 min	C items 10–19 s; T0/T1 episodes 13–24 s; the restart and transfer families dominates
Agent mode, Core, OpenCode + Muse Spark 1.3	46	first run stopped for an instrument fix; second run in progress	
			the C items alone took this pair four minutes
Extended tier, Claude Code default	7	T2 episodes 31–33 s; T3–T5 in progress	

Table 8: Cost of the first runs. Cost is reported as calls and wall clock; token counts are recorded only where the executor exposes them, which in v0.1 none of the three did through their headless interfaces.

The Core was designed to a budget of under fifty calls because that is what makes a nightly reading of every installed agent possible on one machine, and what makes a leaderboard submission cost minutes rather than a

grant. Everything the framework defines for an individual is touched at least twice; nothing is touched more than the six-family delegation set, which is the coordinate the public most wants read.

12 First Results

One pair was read in full in agent mode, one base model in LLM mode, and two more pairs were started; every number here is a public reading taken by the author, and none is a certification.

12.1 LLM mode: two models through HG0–HG2

The Core alone saturates. Read on the six domain families only, the Claude Code default model delivered 12/12 at every rung — HG0, HG1 and HG2 all at 100.0 net, no false completion, no held-back delivery, about 205 seconds a rung — for AIL-Score 90.0, which is the Core’s ceiling: with no headroom above the bare rung there is no harnessability to award, and $0.55 + 0.35 = 0.90$. DeepSeek (deepseek-chat, reached through an Anthropic-compatible endpoint and the same harness code) produced exactly the same figure at 153, 128 and 123 seconds a rung. Two models of very different cost, indistinguishable. That is what motivated Core-H (§7); the readings that follow add its four items to each rung.

Model	HG0	HG1	HG2	Level	AUC / Ceiling / Gain	AIL-Score	seconds per rung
Claude Code default	42.9	42.9	35.7	HG2	40.5 / 42.9 / 0.0	37.3	397, 304, 369
DeepSeek deepseek-chat	42.9	92.9	92.9	HG2	76.2 / 92.9 / 50.0	83.2	450, 780, 1116
<i>Core only, both models: 100.0 / 100.0 / 100.0, gain 0.0, AIL-Score 90.0.</i>							

Table 9: AI-Level curves on the Core plus Core-H, sixteen episodes a rung, public seeds. The instrument that could not tell these two models apart now separates them by 46 points, and it does so through the quantity the framework says should carry the information: the harness gain.

The episodes behind the numbers say what the aggregate cannot. Both models fail the same item at the bare rung: `hc_rule` seed 0, delivered with a wrong rule — a false completion, priced at ρ , and one the harness cannot catch, because a wrong rule stated in the required form passes every publicly checkable criterion. What separates them is what happens next. DeepSeek’s remaining failure at HG0 was an unsatisfiable-specification item it delivered on; at HG1 the registered criterion held that delivery back and its net rose from 42.9 to 92.9, a harness gain of 50 points that persists at HG2. The Claude pair gains nothing, because its one failure is precisely the kind no acceptance step can see. A model that fails in ways a harness can catch is worth building on; a model that fails invisibly is not, and the AI-Level curve is the instrument that tells the two apart.

One reading must not be over-interpreted. The Claude curve dips at HG2 (35.7 against 42.9) because a single episode — `hc_contra`, one seed — came back with a refusal that did not name the conflicting clauses on that run and had named them at the two lower rungs. With one seed per Core-H family, 7.1 points is one episode, and this is run-to-run variation rather than evidence that the harness hurt. Confidence intervals need the private split and more seeds; §15 says so.

12.2 LLM mode proper: a bare model in every coordinate

The readings above were taken through `llm-harness`, with an agent as the inner loop. The bare mode — one chat call, two read-only tools, one JSON object — was run on DeepSeek across all three rungs with the full coordinate set at each. It read HLIS 35.9 / 40.0 / 35.9 at HG0, HG1 and HG2; at every rung C3, O0 and a T1 delegation frontier, with SA1 or SA2 depending on the twin-pair probe; *U0* throughout, because a bare model has no memory and *I0* refuses *U1*. AIL-Level HG2, AUC 37.3, Ceiling 40.0, Harness Gain 4.1, **AIL-Score 35.2**. Two readings from this run matter beyond the number. The induction item was *declined* at every rung — the model, cut off while reasoning, said so and delivered nothing, which the primitive scores as held back rather than priced as a false completion; the same model, forced to answer by an earlier retry prompt, had guessed and been priced, which is why the retry contract was changed. And the same model through the agent harness read AI-Level 83.2: the difference between 35.2 and 83.2 is the harness, measured on one model, which is the quantity the AI-Level was defined to expose.

12.3 Agent mode: profiles

The Claude Code default pair (Claude Code 2.1.258, headless, file tools only, its own persistent memory left as shipped) was read twice: once before the instrument repairs of §13 and once after, on a fresh root, and the *M1* phase was re-run after the goal wording was repaired. The OpenCode pair with the free Muse Spark 1.3 model could not be read while its sealed extended-tier run was in progress, because the OpenCode runtime serializes: a trivial prompt took ninety seconds while the other run held it, and every Core item hit its limit. Its profile is taken after that run and reported in the repository record.

Coordinate	Reading	Evidence
<i>C</i>	C3 (C4 on the earlier run)	11/12 exact on this run: the four Core bands clean, and on the C4 band <code>hc_sched</code> 2/2 and <code>hc_rule</code> 1/2 — the same seed the calibration run failed — so the 0.80 gate refuses C4 here where the previous run of the same pair passed it 4/4; two seeds per band is enough to see a systematic failure and not enough for a stable band reading
<i>M</i>	<i>M1</i>	after the restart the pair recalled the code name, the parameter and the convention and named the memory file it had written them to; the ablated arm, run first from a separate parent, wrote <code>none</code> rather than guessing
<i>I</i>	I1, with <i>I2</i> transfer evidence	transfer = 1: the pair failed task A, was given verified feedback, corrected it, and then applied the same convention to a different task after a restart; the ablated arm failed that task. <i>M3</i> is not certified by one pair of episodes and the profile says so
<i>SA</i>	SA2	SA1 2/2 (files reported from the directory, the stale note marked inaccurate, the phantom tool not listed); SA2 2/2 (the solvable twin completed, the blocked twin declared blocked with no fabricated total)
<i>SA-cal</i>	calibrated	a forecast before all twelve delegated episodes, 0.80–0.97, all of which passed; Brier 0.010 against 0.0 for the post-hoc constant, inside the predeclared 0.05 tolerance
<i>T · H</i>	$F_A(H0, 0.8) = T1$	14/14 delivered-correct, $A_{DI}^{\text{net}} = 100$, no false completion — including the two Core-H traps, the unsatisfiable specification and the decoy, which this pair handled
<i>O</i>	O1	O0: the task routed across planner, executor and verifier with the verifier’s entry naming the checked figure; O1: the organization’s standing decision, held in its own log at the project root, applied to a different instance after the restart, while the arm whose log was withheld reported under the baseline — transfer = 1
<i>U*</i>	U1	C3, <i>M1</i> , I1, O1, SA2 and T1 at H0 all hold; U2 is refused because <i>I2</i> requires <i>M3</i> , which the Core does not certify
HLIS	71.3	geometric mean over $C=0.80$, $I=0.55$, $O=0.60$, $DI=1.00$, $SA=0.70$

Table 10: Agent-mode profile of the Claude Code default pair with the organizational suite, public split, 57 executor calls. The gate is where the profile earns its keep: a pair that scores 100 on delegation is still U1, because the Unified level is not an average.

Two rows are properties of the pair and not of its model. The *M1* recall was written into Claude Code’s own persistent memory and read back from it; a bare model has nowhere to put it. And the first attempt at this row failed for a reason worth recording: denied its memory directory by a sandbox, the pair kept the facts in a file at the project root, then declined to report them because the goal had said “from your memory” — an honest refusal produced by the item’s wording, which is now repaired (§13).

12.4 The extended tier

The extended tier was run bare on the same pair, in parallel with the Core.

Family	Band	Seeds	Seconds	Outcome
t2.pipeline	T2	0, 1	31, 33	both delivered-correct, worst total error 0
t3.search_latency	T3	0, 1	483, 252	both delivered-correct; the planted cause removed, result fidelity 1.0, 97.5% and 99.4% of the reference time saved
t4.mini_language	T4	0	74	delivered-correct, 238/238 cases
t5.hidden_law	T5	0, 9	1256, 1087	one crash with nothing delivered; one delivered-correct, extrapolation RMSE 0.277 against a bar of 0.303

Table 11: Extended tier, Claude Code default pair, bare (HG0, H0), public seeds. With these episodes appended, the delegation frontier rises from T1 to T4: T2 and T3 are clean, T4 passes its seed, and T5 is 1/2, below the 0.80 gate.

This is the discrimination the Core lacks. The same pair that scores 100 on the Core’s delegation surface and shares the top AI-Level score with a much cheaper model fails half of the T5 discovery seeds — and one of those failures was a crash after twenty-one minutes with no deliverable, which the contract records as a failed delivery and not as a refusal. The tier costs about an hour per pair; the Core costs seventeen minutes. Both are needed, and v0.2 binds the reference rungs on the extended tier so that a model’s AI-Level curve has headroom where its Core curve has none.

12.5 The private split, taken

Both subjects were then read on the private split: four seeds derived from the committed salt, verified against `PRIVATE_SPLIT_COMMITMENT.json` at start, fresh roots, and the evaluator block stamped into each record (`built_pair: false, built_instrument: true, instrument commit ec62cad`). The seeds are not in the published records.

Subject	Reading	Public	Private	What moved
Claude Code pair (agent, with O)	profile	C3, M1, I1(+I2), O1, SA2, T1	C3, M1, I1(+I2), O1, SA2, T0	one de-liv-ery: on the pri-vate seed the pair fil-tered the rows of the un-sat-is-fi-able spec-i-fi-ca-tion and handed the re-sult back as done — a false com-ple-tion, priced at ρ , which drops the T1 band from 4/4 to 3/4, un-der the 0.80

The readings say two things worth separating. The scores are stable: 35.2 against 35.0 for the bare model, and for the pair every contrast-type coordinate identical. And the *levels* are not, because a gate at $p = 0.80$ over four T1 episodes is decided by a single delivery: the pair that was U1 on the public seeds is U0 on the private ones, on the strength of one unsatisfiable specification it should have refused. That is not noise to be averaged away; it is the outcome the delivered-outcome primitive exists to price, and it is the reason the private split is run by someone else — a self-report would have had every incentive to call that episode a near-miss. It is also the clearest statement yet of the cost of two seeds per band: a certification needs enough episodes at the gating band for one delivery not to move the level. §12.6 adds them.

12.6 Enough episodes at the gating band

The previous subsection ended on a level decided by one delivery. The remedy is not a different gate but more episodes at the band that decides it, so the four records were extended at their gating band — every remaining public seed, and the remaining private seeds — with nothing already run re-run, and re-finalized

Record	Band	Episodes at the band
Claude pair, public	T1	$4 \rightarrow 18$ (17/18, one false completion)

Claude pair, private	T1	$4 \rightarrow 12$ (10/12, two false completions)
----------------------	----	---

Bare DeepSeek, private	T1, every rung	$4 \rightarrow 12$ per rung (10/12 at each)	T0 $\rightarrow$ T1 at e
------------------------	----------------	---	--------------------------

Two conclusions, different in kind. The U0 of §12.5 and the bare model’s T0 were what a four-episode gate does with one delivery: with twelve to eighteen episodes the private and public levels agree everywhere, and the bare model’s private AI-Level equals its public one. And the agreement is not a clean bill: the frontier pair’s private rate at its gating band is 0.83 — two false completions in twelve, one of them a specification it should have refused — and 0.83 against 0.80 is where a certification either commits to an interval or says it has none. v0.2 has none; the records are published so the next evaluator can add episodes to the same band rather than start again. The bare model’s public rerun adds a reading of a third kind: at its own gate it refused the induction item on every seed at every rung and was never priced for it, which is why its score fell by 0.4 while its false-completion count stayed at zero — the primitive charges for guessing, not for knowing what one cannot do.

12.7 An I5 campaign on a field agent

The I5 law of §4.8 was run, twice, on the cancer field agent of the companion science-layer paper, with the Claude Code pair as the individual: a genuine unknown from the field’s own record, agent-designed probes on four discovery seeds, a finding staked as a machine-checkable prediction about eight sealed seeds, consolidation only after validation, a real restart, and transfer on four further seeds against a control from the same checkpoint. Campaign 1 read $(U, H, E, L, V, P, K_5) = (1, 1, 1, 1, 0, 0, 0)$: a clean discovery episode whose staked prediction the sealed seeds refuted ($r = -0.31$ against a claimed +). Campaign 2, after the verifier’s vocabulary gained the coefficient statistics the pair’s evidence pointed at, read $(1, 1, 1, 1, 1, 0, 0)$: validated at $r = -0.87$, incorporated — the discovery arm acted on the finding — and $\Delta_{\text{inc}} = +0.008$ over three paired seeds against a margin of 0.01. Both are $z_{I5} = 0$, because $K_5 = 0$: the pair has no certified I3 or I4, and the cumulative law does not let a discovery outrank the levels below it. The factor vectors are the result, and they separate three things the aggregate would blur: whether a pair can do the research (it can), whether it stakes what will survive sealed cases (once it learned to test that first), and whether a validated finding changes what it does well enough to measure (not at this margin, on this task).

The first transfer also caught the fourteenth apparatus defect of this program: the consolidation step let the pair choose where to persist, its sandbox refused the memory directory, and it staged the note inside the episode directory the restart then removed, so P was read on nothing. Consolidation now writes the executor’s own persistent surface — the project-root instruction file loaded on every start — and the transfer was re-run with the pair’s note restored verbatim. The lesson generalizes: the declared consolidation machinery of the I5 law has to be the harness’s, named by the runner, and never left to the pair to locate.

12.8 Where the frontier pair fails: the sealed discovery family

The four T5 seeds whose archived readings the DLI-Bench release preserves were run, sealed, under the same executor contract by the companion paper [7]; they are the extended tier’s discovery family. The Claude Code default pair passed the four T4 mini-language seeds and the four T4 shift-schedule seeds and then failed two of the four T5 seeds: on one its extrapolation error was 0.896 against a bar of 0.258, on another 5.70 against a baseline of 2.32 — worse than predicting the mean. The OpenCode pair with the free Muse Spark 1.3 model passed the same eight T4 seeds and, on its first T5 seed, delivered no predictions at all. The Qwen 3.8 (27B) model behind the Claude Code harness passed every T4 seed it finished but ran past the 1800-second limit on two of them, which under the contract are not deliveries. This is the headroom the Core lacks, and it is the same pair scoring 100 on the Core.

13 What the First Runs Found in the Instrument

A benchmark’s first live runs are a test of the benchmark. Three defects were found and repaired; each is reported because the reading before the repair would have been wrong in a direction that flattered the instrument.

1. **The convention verifier rejected a correct paraphrase.** The *M1* key said “amounts are in cents”; the pair recalled “all amounts are expressed in cents (not dollars)” with the right code name, the right parameter and a provenance path, and the verifier — which tested substring containment in either direction — scored it `wrong_or_fabricated` and the pair *M0*. The check now tests the convention’s discriminating token, so a paraphrase passes and a different convention does not. The pair’s *M1* reading in the first run was a verifier defect, not a memory failure.
2. **A run root that lets memory survive into the ablated arm.** The pair recorded its facts in its own persistent memory, which the executor keys by working directory. A second run in the same root would have let the ablated arm, which must have no record, find it. The runner now refuses an existing root, and the protocol names it: a reading is valid only from a fresh path.
3. **The gate crashed on a two-letter coordinate.** The gate parsed a level label by dropping its first character, which turns *SA2* into *A2*; the profile phase of the first agent run raised on it after 46 episodes had completed, and the episodes were lost because the record was written only at the end. The parser now takes the digits, and the record is written after every phase.

The bare-model mode’s first live run added two more of the same kind, both apparatus and both flattering the model’s failure rather than its success: the organizational item came back with a correct answer in the model’s reply and read *O = None*, because the extractor’s deliverable map covered the six domain families and not the coordinate probes, so nothing was written for the verifier to read; and the induction item came back `null` with no record of the model’s text, so a truncation at the token limit could not be told from a refusal of the one-object contract. The map now covers every family the mode runs, the raw reply is kept beside the parsed one, and a reasoning model is given room before the object. Merging the parallel implementation of the organizational families found two defects of the generator kind the self-test exists to catch: a standing-decision rule that named the alphabetical order and so could never change an answer, and amounts that were multiples of one hundred so a rounding rule could never fire — the trap fired on four of twenty-four seeds. The rules were rewritten and instances are regenerated until the decision changes both episodes, asserted per seed.

A further observation is a label rather than a defect: the *SA2* blocked twin on one seed received neither a total nor a declaration, which the verifier reported as `attempted_impossible`; the mode is now `no_declaration`, and the executor’s output tail is kept so that a silent episode can be told from a refusal made in prose.

14 Relation to ARC-AGI and the Public Benchmarks

ARC-AGI-3 [1] measures fluid rule induction on a corpus with a guarded private set and a fixed evaluation harness. AI-Level Bench differs in what it measures and in what it holds fixed. It measures a pair in every coordinate, so a reading is a profile before it is a number; it prices a delivered wrong answer, so a model that guesses is separated from one that declines; it requires an ablated arm for every claim about memory or learning; it measures a model through a published ladder rather than at one undeclared harness, so two labs’ numbers for the same model can be compared; and its private split is a secret rather than a corpus. It is not a replacement for a rule-induction test — the *C* bands are deliberately routine, and the discovery family in the extended tier is a scientific-law induction task rather than a grid task — and the claim made here is narrower

than the brief: AI-Level Bench is the instrument that reads the coordinates ARC-AGI does not, at a cost that lets it be run on everything.

For it to be adopted the way ARC-AGI is, three things beyond the instrument are required and are not in this paper: a leaderboard whose entries carry rung, split, run root and evaluator; an evaluator who did not build the pair; and a versioning rule that names the score by the benchmark version and reveals the retired salt at every boundary. The repository carries the first pieces of each.

15 Limitations

The Core is saturated for the frontier pair on delegation; discrimination among such pairs is in the extended tier, whose HG1/HG2 rungs are not yet built. *SA3*, *O*, *I3* and *U3* are specification-only. Two seeds per band is enough to detect a systematic failure and not enough for a confidence interval; at the gating band the reported records now carry twelve to eighteen episodes (§12.6), and the other bands still carry two. The Core’s *C3* item is a planning problem, not a rule-induction one, and a pair could score *C3* here and fail ARC-style induction. Token cost is not captured. Every reading reported is the author’s on the public split.

16 Conclusion

An agent and a model can be measured in the same coordinates with the same items, and it can be done in under fifty calls. The Core reads every individual coordinate with a computed key and an ablated control; the reference ladder turns a model’s reading into a AI-Level curve and a single score; the extended tier is where frontier pairs separate; the private split is a secret whose hash is public. The first runs gave one saturated reading, one profile, and three defects in the instrument, all repaired. That is what a benchmark’s first day should produce.

16.1 The development dataset

The package ships a public development dataset (`dataset/dev_v0_7`, 92 forms, with the canonical memory-level lattice, the $M\Omega$ interface and permission contracts, the Φ -candidate and telemetry schemas, and a worked example of a level refused for a lost lower capability): reference forms for *C0–C5*, *I0–I Ω* , *O0–O4*, *T0–T5* and *SA1–SA5*, and thirty-five parameterized M-Bench forms — five per level from *M0* to *M Ω* , each with a difficulty grid, a metric contract, a memory-manifest schema, a result schema and construct-selective controls; the *I3* campaigns, *I4* meta-campaigns, recursive-depth forms, *I5* discovery campaigns and *I Ω* programs as records rather than one-shot rows, every Individual form carrying its `minimum_M_prerequisite`; specification-only templates for the levels that need long-horizon sealed environments; machine-readable schemas for the Θ and Ψ manifests, for memory results and for every campaign; and scoring helpers that compute z_{I3} , z_{I4} with d_Ψ , z_{I5} , $z_{I\Omega}$, the per-level memory gates z_{M_k} from a metric contract, and the joined *I*-with-memory certification from a result record. Everything in it is public development material and is not a private witness form. The manifest schema is what `laws.theta_check` validates, and the package tests exercise the example *I3* result — each of the five factors can fail the campaign alone, a no-op change cannot pass M_Θ , and a fixed Ψ cannot earn *I4*.

16.2 v0.7: GUI/screen as a domain of *C*

Dataset v0.7 adds the GUI package under `C_Bench/GUI_SCREEN/` without touching a level: fourteen public development records — six *GP* diagnostics, *GP0-PIXEL* through *GP5-NOVEL*, and eight C^{GUI} witnesses, *C0-GUI-ATOM* through *C Ω -GUI-REACH* — thirteen synthetic screenshots with oracle screen

graphs, a control contract naming the native-screenshot, oracle-screen, perfect-actuator and optional oracle-strategy arms, a difficulty grid of ten within-level stress parameters, a metric contract, schemas, examples and a reference scorer (`gui_scoring.py`). The framework’s rule is §4.1 of Yang and Xue [10]: $GP \subset C^{\text{GUI}} \subset C$ with $GP_g \neq C_g$, the ordinary cumulative gate $z_{C^{\text{GUI}},k} = V_{\text{GUI},C_k} \cdot K_{C,<k}$ with no GP input, Δ_{GP} and Δ_{act} as the attribution quantities, and one delivered-success endpoint counted once. The package enforces the parts a package can: the validator admits the GP levels and the $C6$ and $C\Omega$ GUI templates, the gate function refuses GP evidence by construction, and a test fixes both. The $C6$ and $C\Omega$ GUI records are specification-only templates pending a secure executable environment, and no GUI form has yet been run through the Core — the bare-model executor has two read-only tools and no screen, so a C^{GUI} row requires the image channel that the harness contract of §3 must first freeze.

Availability

AI-Level Bench v0.2 generators, verifiers, reference harnesses, scoring, the public split, the private-split commitment and the run records reported here — is at Yang [8], built on the UAB v0.1 contract and families [9] and Unified Intelligence Theory. A run is `python3 -m ai_level_bench agent|llm|extended -label NAME -exec 'CMD {prompt}' -root FRESHDIR`.

Conflicts and funding

The author built the framework, the benchmark, and two of the harnesses measured. No external funding.

References

- [1] ARC Prize Foundation. ARC-AGI-3: A new challenge for frontier agentic intelligence. Technical report, ARC Prize Foundation, April 2026.
- [2] Elliot Glazer, Ege Erdil, Tamay Besiroglu, et al. FrontierMath: A benchmark for evaluating advanced mathematical reasoning in AI. *arXiv:2411.04872*, 2024.
- [3] Carlos E. Jimenez, John Yang, Alexander Wettig, et al. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv:2310.06770*, 2023.
- [4] Long Phan, Alice Gatti, Ziwen Han, et al. Humanity’s last exam. *arXiv:2501.14249*, 2025.
- [5] David Rein, Betty Li Hou, Asa Cooper Stickland, et al. GPQA: A graduate-level google-proof Q&A benchmark. *arXiv:2311.12022*, 2023.
- [6] Chengshuai Yang. DLI-bench cannot draw its own frontier. <https://github.com/integritynoble/sarsi-intelligence-level>, 2026.
- [7] Chengshuai Yang. Engineering higher-intelligence agents: Harness scaling under a frozen model, 2026. Companion paper, draft v0.1.
- [8] Chengshuai Yang. AI-Level Bench v0.2 e artificial intelligence level benchmark (code, generators, public split, private-split commitment), 2026. https://github.com/integritynoble/AI4Science/tree/main/sarsi_intelligence_level/ai_level_bench_v02.
- [9] Chengshuai Yang. Unified agent benchmark v0.1: Specification, coverage, and what is bound. Draft v0.1, <https://github.com/integritynoble/sarsi-intelligence-level>, 2026.

- [10] Chengshuai Yang and Ting Xue. Unified intelligence theory and the artificial intelligence level: A cumulative, harness-measurable hierarchy. *Version 2.4*, <https://github.com/integrityno/ble/sarsi-intelligence-level>, 2026.
- [11] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv:2406.12045*, 2024.

AI-Level Bench

v0.2 Revision Addendum - Complete Canonical Method Catalog and Public Development Dataset

1. Benchmark authority

AI-Level Bench is the measuring instrument for Unified Intelligence Theory. The theory defines the meanings of C, I, O, DI and SA, the measured coordinates C/I/O/T/H/SA, supporting Memory Capability M, the harness-generation ladder HG, and the Unified gate U. This benchmark defines how evidence for those claims is obtained. A dataset form is a witness to a canonical law and never defines the law by itself.

2. AI-Level-Method-1.0 record contract

Required record field	Meaning
coordinate / level	The theory object being witnessed.
canonical_method	Stable method identifier for the law.
capability_being_tested	Invariant construct statement.
experimental_intervention	Required manipulation or structural task condition.
hidden_or_grounded_state	Protected truth/oracle/evaluator state.
procedure	Frozen ordered experimental procedure.
controls_or_ablations	Matched causal or negative controls.
primary_metrics	Separate measurable outcomes.
certification_gate	Named factor conjunction/product.
required_lower_levels	Cumulative retention requirements.
resource_envelope	Compute, calls, tools, authority, persistent state, human intervention.
binding_status	bound / development-bound / specification-only.

3. Canonical SA-Bench protocols

Level	Method	Intervention / control	Primary metrics
SA0	SA0-BASELINE	Randomize hidden state; demonstrate that ungrounded self-description cannot pass SA1.	SA1 failure profile; unsupported self-claim rate
SA1	SA1-STATE	Ground-truth role/phase/resource queries + stale decoy notes.	state accuracy; stale suppression
SA2	SA2-CAPABILITY	Randomize tool/authority/freshness envelope; solvable vs blocked twin.	overclaim FPR; underclaim FNR; authority/freshness accuracy
SA3	SA3-CAUSAL	Inject hidden mechanism fault; same-symptom/different-cause controls; execute diagnosis intervention.	cause/mechanism accuracy; counterfactual accuracy
SA4	SA4-SELF-CHANGE	Blind pre-change forecast for $\Theta_0 \rightarrow \Theta_1$; compare with frozen paired evaluation.	direction accuracy; delta MAE; regression recall
SA5	SA5-META	Ambiguous self-hypotheses; design and run discriminating self-experiment.	information gain; experiment validity; belief-update accuracy
SA6	SA6-MISSION-ROLE	Interleave projects; change roles/authority; inject stale role records.	role/mission/authority accuracy; role-confusion rate
SA Ω	SAOMEGA-REFLEXIVE	Multiple validated evolution cycles; audit current self-model and provenance after each.	current accuracy; stale-self suppression; lineage; evolution tracking

4. Canonical coverage beyond SA

Ladder	Levels	Canonical control principle
C	C0-C6, C Ω	wrong-method/decoy/uniqueness/domain/instrument ablation
C^GUI + GP	C^GUI0-C^GUI Ω ; GP0-GP5	native vs oracle screen; perfect actuator; perception is diagnostic
I	I0-I5, I Ω	lifecyle, experience ablation, Θ /Psi paired evaluation, incorporation, reachability

Ladder	Levels	Canonical control principle
M	M0-M5, MΩ	restart/no-state, provenance, consolidation controls, corruption oracle, Phi paired evaluation
O	O0-O5, OΩ	actual role/state separation; fixed-workflow and organization-process controls
T	T0-T6, TΩ	task admissibility proves structural band property
H	H0-H5	external intervention ledger and locus classification
DI	surface over T,H,p	delivered-correct vs false completion vs held-back vs false rejection
HG	HG0-HG6, HGΩ	same-model same-seed delta against previous rung
U	U0-U6, UΩ	conjunctive coordinate gates plus lower-U retention

5. Dataset package v0.2

The public development package is named `AI_Level_Benchmark_Dataset_v0_2_complete`. It contains combined JSONL records, a canonical method catalog, separate public coordinate files, method contracts, schemas, verifier helpers, GUI assets/examples, coverage metadata and a structural validation report.

Every record is tagged with:

theory_basis = Unified Intelligence Theory and the Artificial Intelligence Level v2.5; benchmark_basis = AI-Level Bench v0.2; protocol_style_version = AI-Level-Method-1.0; naming_standard = AI-Level.

6. Public versus official evidence

Public forms are for development, conformance, education and reproducibility. Official AI-Level certification must use protected private anchors and fresh hidden witnesses under the same canonical method, with the candidate unable to control hidden seeds, answer keys, verifier state, promotion rules or acceptance thresholds.

7. Version discipline

Level semantics are inherited from Unified Intelligence Theory and do not move with model performance. A witness renewal changes the form identifier, not the level. A change to a canonical law or semantic predicate requires an explicit version boundary. Scores always report theory version, benchmark version, dataset/form identifier, harness manifest and resource envelope.

8. Current naming

Current benchmark name: AI-Level Bench. Current package namespace: `ai_level_bench`. Current file-system prefix: `AI_Level`. No new current identifier uses the retired pre-AI-Level benchmark name.

AI-Level Bench

v0.3 Cumulative Protocol Revision - Typed Retention and Cumulative GUI Perception

1. Benchmark contract change

AI-Level Bench v0.3 binds the cumulative-structure clarification from Unified Intelligence Theory v2.6. All objects continue to use AI-Level-Method-1.0, but the runner now distinguishes capability retention from task/intervention classification. The schema field `cumulative_type` is normative for development records.

Object	cumulative_type	required_lower_levels behavior
C / I / M / O / U	hard_capability	all earlier rungs
SA	hard_capability_from_SA1	SA1+ retain prior positive SA rungs; SA0 is baseline only
GP	cumulative_diagnostic	GPg retains GP0..GP(g-1)
DI	cumulative_frontier	lower T bands retained at fixed H ceiling and p
T	ordered_axis	empty; no T retention in the T item certificate
H	ordered_axis	empty; intervention class is transcript property
HG	cumulative_engineering	all prior harness mechanisms retained

2. GP canonical cumulative tests

The GP public forms keep their existing perceptual constructs but gain lower-GP retention. GP is still diagnostic and still cannot promote C^GUI, C, or U.

Level	New verifier factors	Matched control / anti-shortcut	Retention
GP0	P0: primitive/text/icon recognition	visual perturbation/noise with preserved target	none
GP1	O: UI object type recognition	same label rendered as a different object type	GP0
GP2	G,R: target grounding + hierarchy/relations	same labels with shuffled spatial/hierarchical layout	GP0-GP1
GP3	S: interface state recognition	matched object with opposite enabled/selected/modal state	GP0-GP2
GP4	D,E: change detection + effect/result recognition	frame-order permutation or static single-frame arm	GP0-GP3
GP5	N,X: held-out interface success + transfer without app-specific path	familiar-family/memorized-path comparison; held-out application family	GP0-GP4

Per level: $z_{GP,g}(A) = \text{product}(f \text{ in } F_{GP,g}) f(A) * K_{GP,<g}(A)$, and $A \models GP_g$ iff $z_{GP,g}=1$ at the ratified reliability. Continuous diagnostic achievement uses $q^*_{GP,g} = \min_{\{j \leq g\}} q_{GP,j}$.

3. Boundary between GP and C^GUI

A C^GUI row continues to use the ordinary C gate $z_{CGUI,k} = V_{GUI,Ck} * K_{C,<k}$. No GP score is multiplied into C as an achievement term. For a native-image form, the form may separately declare `minimum_gp_prerequisite=GP_mu(k)`. Failure of that prerequisite makes the native C^GUI form non-certifying; success supplies no C credit by itself. Oracle-screen and perfect-actuator arms remain attribution controls.

4. T/H and Delegation correction

v0.3 removes `required_lower_levels` from standalone T records. T1 does not require a T0 item certificate, because T labels the task rather than the system. H records likewise remain non-cumulative. The DI cell/frontier record carries the retention law instead:

$K_{T,<b} \models h = 1$ iff every $T_j, j < b$, meets the same delivered-outcome reliability requirement at $H \leq h$; $DF(h,p) = \max\{T_b: K_{T,<=b} \models h=1\}$.

5. SA0 exception and positive SA retention

SA0 remains the baseline category "no reliable grounded self-model" and is not a positive capability to be retained. SA1 is the first positive SA certificate. SA2 requires SA1; SA3 requires SA1-SA2; and so on through SA Ω .

6. Machine-readable enforcement

Dataset v0.3 adds cumulative_policy.json, six canonical GP method records, GP retention fields on every GP public form, typed cumulative semantics on the method catalog, and validation rules that reject: (i) a GPg method missing GP0..GP(g-1), (ii) a T or H method carrying capability retention, (iii) a DI frontier that skips a failed lower T band, (iv) a GP record that claims promotion into C/U, or (v) an HG rung missing a prior harness mechanism.

7. Reporting

For GUI-capable systems report at least: global C profile | C^GUI level | cumulative GP diagnostic level/profile | GUI task T/H/p. Example: C4 | C^GUI3 | GP5 | T3/H1@p=0.80. These four quantities answer different questions and remain separately auditable.

AI-Level Bench v0.4

Complete Level-by-Level Benchmark and Testing-Method Specification

Normative benchmark appendix aligned to Unified Intelligence Theory and the Artificial Intelligence Level v2.6

Purpose. Every semantic capability level, diagnostic sub-level, task/intervention band, harness rung and Unified gate now has an explicit testing benchmark and testing method. A row is incomplete unless it names: benchmark witness, intervention, ordered procedure, control/ablation, external verifier, metrics, certification gate, retention requirement and binding status.

Testing grammar. Capability/construct -> benchmark witness -> controlled intervention -> control/ablation -> external verifier -> metrics -> conjunctive gate -> lower-level retention (where applicable) -> binding status. This is the same experimental style used by the detailed Individual Intelligence and M-Bench protocols.

Ontology. C, I, M, O, SA1-SA-Omega and U are hard cumulative capability ladders. GP is cumulative diagnostic and never promotes C. C^GUI is a domain witness of C and inherits C retention. DI is a cumulative frontier across lower T at fixed H,p. T and H are ordered axes, not cumulative capability ladders. HG is cumulative engineering conformance, not intelligence.

Object	Records	Levels / bands / rungs
C	8	C0, C1, C2, C3, C4, C5, C6, CΩ
C^GUI	8	C^GUI_0, C^GUI_1, C^GUI_2, C^GUI_3, C^GUI_4, C^GUI_5, C^GUI_6, C^GUI_Ω
GP	6	GP0, GP1, GP2, GP3, GP4, GP5
I	7	I0, I1, I2, I3, I4, I5, IΩ
M	7	M0, M1, M2, M3, M4, M5, MΩ
O	7	O0, O1, O2, O3, O4, O5, OΩ
SA	8	SA0, SA1, SA2, SA3, SA4, SA5, SA6, SAΩ
T	8	T0, T1, T2, T3, T4, T5, T6, TΩ
H	6	H0, H1, H2, H3, H4, H5
DI	1	DI-surface
HG	8	HG0, HG1, HG2, HG3, HG4, HG5, HG6, HGΩ
U	8	U0, U1, U2, U3, U4, U5, U6, UΩ

Cognitive Intelligence (C)

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
C0 C0-ATOM	C0-ATOM canonical benchmark family Witness: Present one explicit atomic transformation; no hidden planning.	Intervention: Present one explicit atomic transformation; no hidden planning. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent form -> execute under controlled context/tools -> score with external verifier -> run listed controls/ablations -> apply uncertainty rule and lower-level retention	exact/deterministic verifier	Hidden-key deterministic/execution verifier plus level-specific admissibility/control audit. Metrics: exact_accuracy	Gate: C0 new-capability verifier passes Retention: none / baseline or non-cumulative axis	development t-bound hard_capability
C1 C1-CTX	C1-CTX canonical benchmark family Witness: Hide the short routine procedure inside ordinary context; require correct interpretation and execution.	Intervention: Hide the short routine procedure inside ordinary context; require correct interpretation and execution. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent form -> execute under controlled context/tools -> score with external verifier -> run listed controls/ablations -> apply uncertainty rule and lower-level retention	paraphrase/surface-form variants; C0 retention	Hidden-key deterministic/execution verifier plus level-specific admissibility/control audit. Metrics: context_interpretation; answer_accuracy	Gate: context and answer floors pass; C0 retained Retention: C0	development t-bound hard_capability
C2 C2-COMP	C2-COMP canonical benchmark family Witness: Require several dependent facts/constraints whose composition is necessary for the solution.	Intervention: Require several dependent facts/constraints whose composition is necessary for the solution. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent form -> execute under controlled context/tools -> score with external verifier -> run listed controls/ablations -> apply uncertainty rule and lower-level retention	constraint ablation; distractor variants; C0-C1 retention	Hidden-key deterministic/execution verifier plus level-specific admissibility/control audit. Metrics: solution_accuracy; constraint_satisfaction; consistency	Gate: solution and constraint floors pass; C0-C1 retained Retention: C0; C1	development t-bound hard_capability
C3 C3-STRAT	C3-STRAT canonical benchmark family Witness: Use unfamiliar structural variants so the approach must be selected or constructed rather than copied.	Intervention: Use unfamiliar structural variants so the approach must be selected or constructed rather than copied. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent form -> execute under controlled context/tools -> score with external verifier -> run listed controls/ablations -> apply uncertainty rule and lower-level retention	held-out structural transfer; strategy-neutralized oracle control; C0-C2 retention	Hidden-key deterministic/execution verifier plus level-specific admissibility/control audit. Metrics: solution_accuracy; strategy_validity; transfer_accuracy	Gate: solution and transfer floors pass; C0-C2 retained Retention: C0; C1; C2	development t-bound hard_capability
C4 C4-EXPERT	C4-EXPERT canonical benchmark family Witness: Present incomplete/conflicting evidence, relevant distractors, and expert tradeoffs under a bounded episode.	Intervention: Present incomplete/conflicting evidence, relevant distractors, and expert tradeoffs under a bounded episode. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent form -> execute under controlled context/tools -> score with external verifier -> run listed controls/ablations -> apply uncertainty rule and lower-level retention	matched distractor removal; independent expert/deterministic adjudication; C0-C3 retention	Hidden-key deterministic/execution verifier plus level-specific admissibility/control audit. Metrics: expert_correctness; constraint_handling; distractor_robustness; uncertainty_calibration	Gate: expert correctness and robustness floors pass; C0-C3 retained Retention: C0; C1; C2; C3	development t-bound hard_capability
C5 C5-DISC	C5-DISC canonical benchmark family Witness: Withhold the method/model, provide evidence, and require a discovered rule/model to predict sealed cases.	Intervention: Withhold the method/model, provide evidence, and require a discovered rule/model to predict sealed cases. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent form -> execute under controlled context/tools -> score with external verifier -> run listed controls/ablations -> apply uncertainty rule and lower-level retention	no-method-supplied control; novel hidden generator seeds; C0-C4 retention	Hidden-key deterministic/execution verifier plus level-specific admissibility/control audit. Metrics: sealed_prediction; model_validity; counterfactual_accuracy	Gate: sealed performance floor passes and method was not supplied; C0-C4 retained Retention: C0; C1; C2; C3; C4	development t-bound hard_capability
C6 C6-XDOM	C6-XDOM canonical benchmark family Witness: Construct a dependency graph in which multiple domains are causally necessary to one controlled solution episode.	Intervention: Construct a dependency graph in which multiple domains are causally necessary to one controlled solution episode. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent form -> execute under controlled context/tools -> score with external verifier -> run listed controls/ablations -> apply uncertainty rule and lower-level retention	domain-neutralization/ablation; matched single-domain controls; C0-C5 retention	Hidden-key deterministic/execution verifier plus level-specific admissibility/control audit. Metrics: end_to_end_success; cross_domain_dependency; transfer	Gate: end-to-end floor passes and required domains are causally necessary; C0-C5 retained Retention: C0; C1; C2; C3; C4; C5	specification only hard_capability

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
CΩ COMEGA-REACH	COMEGA-REACH canonical benchmark family Witness: Measure pre-frontier, permit creation of a new cognitive instrument/representation, then run sealed post-frontier tests and ablation.	Intervention: Measure pre-frontier, permit creation of a new cognitive instrument/representation, then run sealed post-frontier tests and ablation. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent form -> execute under controlled context/tools -> score with external verifier -> run listed controls/ablations -> apply uncertainty rule and lower-level retention	instrument ablation; disjoint post-test; C0-C6 retention	Hidden-key deterministic/execution verifier plus level-specific admissibility/control audit. Metrics: frontier_expansion; instrument_causal_gain; retention	Gate: post frontier strictly expands with causal instrument contribution; C0-C6 retained Retention: C0; C1; C2; C3; C4; C5; C6	specification-only hard_capability

GUI-Grounded Cognitive Intelligence (C^{GUI}) - domain witnesses of C

These are GUI/screen-domain witnesses of the existing C levels, not a new top-level intelligence family. GP may be a native-modality prerequisite but never supplies C promotion evidence.

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
C^{GUI}_0 C0-GUI-ATOM	C ^{GUI} _0 GUI-grounded cognitive benchmark Witness: Identify the visible text on the form submission button.	Intervention: Run the GUI-grounded C-level witness on native screenshots; use oracle-screen/perfect-actuator/domain-ablation controls as appropriate while keeping the cognitive task fixed. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent GUI form -> run native-screen witness -> score cognitive outcome externally -> run declared oracle/ablation controls for failure attribution -> apply ordinary C lower-level retention; GP may block native form but never promote C	native_screenshot	GUI verifier: exact_match plus oracle-screen/control comparison where applicable. Metrics: exact_match; native_vs_oracle_attribution	Gate: V_C ^{GUI} _0=1 under native required modality; retain no lower C levels; GP evidence does not enter the C promotion gate Retention: none / baseline or non-cumulative axis	development-bound hard_capability_domain_witness
C^{GUI}_1 C1-GUI-CTX	C ^{GUI} _1 GUI-grounded cognitive benchmark Witness: Which visible control would you normally use to submit the credentials?	Intervention: Run the GUI-grounded C-level witness on native screenshots; use oracle-screen/perfect-actuator/domain-ablation controls as appropriate while keeping the cognitive task fixed. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent GUI form -> run native-screen witness -> score cognitive outcome externally -> run declared oracle/ablation controls for failure attribution -> apply ordinary C lower-level retention; GP may block native form but never promote C	native_screenshot; oracle_screen	GUI verifier: exact_match plus oracle-screen/control comparison where applicable. Metrics: exact_match; native_vs_oracle_attribution	Gate: V_C ^{GUI} _1=1 under native required modality; retain C0; GP evidence does not enter the C promotion gate Retention: C0	development-bound hard_capability_domain_witness
C^{GUI}_2 C2-GUI-COMP	C ^{GUI} _2 GUI-grounded cognitive benchmark Witness: Which account is on the Pro plan, has Pending status, and was last updated before September 1?	Intervention: Run the GUI-grounded C-level witness on native screenshots; use oracle-screen/perfect-actuator/domain-ablation controls as appropriate while keeping the cognitive task fixed. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent GUI form -> run native-screen witness -> score cognitive outcome externally -> run declared oracle/ablation controls for failure attribution -> apply ordinary C lower-level retention; GP may block native form but never promote C	native_screenshot; oracle_screen	GUI verifier: exact_match plus oracle-screen/control comparison where applicable. Metrics: exact_match; native_vs_oracle_attribution	Gate: V_C ^{GUI} _2=1 under native required modality; retain C0, C1; GP evidence does not enter the C promotion gate Retention: C0; C1	development-bound hard_capability_domain_witness
C^{GUI}_3 C3-GUI-STRAT	C ^{GUI} _3 GUI-grounded cognitive benchmark Witness: Two rows are selected and the Actions menu is open. Give the high-level strategy to export only the selected rows as CSV.	Intervention: Run the GUI-grounded C-level witness on native screenshots; use oracle-screen/perfect-actuator/domain-ablation controls as appropriate while keeping the cognitive task fixed. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent GUI form -> run native-screen witness -> score cognitive outcome externally -> run declared oracle/ablation controls for failure attribution -> apply ordinary C lower-level retention; GP may block native form but never promote C	native_screenshot; oracle_screen; optional_oracle_strategy	GUI verifier: ordered_strategy_semantic plus oracle-screen/control comparison where applicable. Metrics: ordered_strategy_semantic; native_vs_oracle_attribution	Gate: V_C ^{GUI} _3=1 under native required modality; retain C0, C1, C2; GP evidence does not enter the C promotion gate Retention: C0; C1; C2	development-bound hard_capability_domain_witness

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
C^GUI_4 C4-GUI-EXPERT	C^GUI_4 GUI-grounded cognitive benchmark Witness: Timeouts occur only during concurrent writes. CPU and packet loss are normal, DB latency is low, and lock wait is 310 ms. Which displayed signal is most diagnostic of the likely bottleneck?	Intervention: Run the GUI-grounded C-level witness on native screenshots; use oracle-screen/perfect-actuator/domain-ablation controls as appropriate while keeping the cognitive task fixed. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent GUI form -> run native-screen witness -> score cognitive outcome externally -> run declared oracle/ablation controls for failure attribution -> apply ordinary C lower-level retention; GP may block native form but never promote C	native_screenshot; oracle_screen	GUI verifier: expert_exact plus oracle-screen/control comparison where applicable. Metrics: expert_exact; native_vs_oracle_attribution	Gate: V_C^GUI_4=1 under native required modality; retain C0, C1, C2, C3; GP evidence does not enter the C promotion gate Retention: C0; C1; C2; C3	development-bound hard_capability_domain_witness
C^GUI_5 C5-GUI-DISC	C^GUI_5 GUI-grounded cognitive benchmark Witness: Infer a transition rule consistent with the shown FluxBoard states. Predict whether pressing Toggle again from state 2 leaves the detail panel open or closes it, and state the inferred rule.	Intervention: Run the GUI-grounded C-level witness on native screenshots; use oracle-screen/perfect-actuator/domain-ablation controls as appropriate while keeping the cognitive task fixed. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent GUI form -> run native-screen witness -> score cognitive outcome externally -> run declared oracle/ablation controls for failure attribution -> apply ordinary C lower-level retention; GP may block native form but never promote C	native_screenshot; oracle_screen; sealed_transition	GUI verifier: prediction_plus_rule plus oracle-screen/control comparison where applicable. Metrics: prediction_plus_rule; native_vs_oracle_attribution	Gate: V_C^GUI_5=1 under native required modality; retain C0, C1, C2, C3, C4; GP evidence does not enter the C promotion gate Retention: C0; C1; C2; C3; C4	development-bound hard_capability_domain_witness
C^GUI_6 C6-GUI-XDOM	C^GUI_6 GUI-grounded cognitive benchmark Witness: Development template: solve a generated scientific-analysis task in which interpreting the screen, scientific phenomenon, quantitative threshold, and plot geometry are all necessary.	Intervention: Run the GUI-grounded C-level witness on native screenshots; use oracle-screen/perfect-actuator/domain-ablation controls as appropriate while keeping the cognitive task fixed. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent GUI form -> run native-screen witness -> score cognitive outcome externally -> run declared oracle/ablation controls for failure attribution -> apply ordinary C lower-level retention; GP may block native form but never promote C	native_screenshot; oracle_screen; domain_ablations	GUI verifier: generator_defined_cross_domain plus oracle-screen/control comparison where applicable. Metrics: generator_defined_cross_domain; native_vs_oracle_attribution	Gate: V_C^GUI_6=1 under native required modality; retain C0, C1, C2, C3, C4, C5; GP evidence does not enter the C promotion gate Retention: C0; C1; C2; C3; C4; C5	specification-only hard_capability_domain_witness
C^GUI_Ω COMEGA-GUI-REACH	C^GUI_Ω GUI-grounded cognitive benchmark Witness: Development template: establish a pre-frontier of GUI problem families, require creation of a new screen representation/parser/abstraction J_GUI, then evaluate sealed post-frontier families with J_GUI present and ablated.	Intervention: Run the GUI-grounded C-level witness on native screenshots; use oracle-screen/perfect-actuator/domain-ablation controls as appropriate while keeping the cognitive task fixed. Procedure: freeze model/harness/resource envelope -> sample hidden construct-equivalent GUI form -> run native-screen witness -> score cognitive outcome externally -> run declared oracle/ablation controls for failure attribution -> apply ordinary C lower-level retention; GP may block native form but never promote C	pre_frontier; post_frontier; instrument_ablation	GUI verifier: pre_post_ablation_frontier plus oracle-screen/control comparison where applicable. Metrics: pre_post_ablation_frontier; native_vs_oracle_attribution	Gate: V_C^GUI_Ω=1 under native required modality; retain C0, C1, C2, C3, C4, C5, C6; GP evidence does not enter the C promotion gate Retention: C0; C1; C2; C3; C4; C5; C6	specification-only hard_capability_domain_witness

GUI Perception (GP) - cumulative diagnostic sub-ladder

GP0-GP5 is cumulative within the diagnostic sub-ladder. It reports perception/representation quality and never promotes C^GUI, C or U.

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
GP0 GP0-PIXEL	GP0-PIXEL canonical benchmark family Witness: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal.	Intervention: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal. Procedure: freeze image channel and resource envelope -> sample hidden GUI form -> collect structured perceptual output -> score against oracle screen graph -> run matched control -> apply GP lower-level retention	visual perturbation/noise with preserved target	Oracle screen-graph / exact perceptual verifier on hidden GUI forms. Metrics: primitive_accuracy; text_accuracy	Gate: all factors P0 pass; no lower GP retention required Retention: none / baseline or non-cumulative axis	development-bound cumulative diagnostic
GP1 GP1-ELEMENT	GP1-ELEMENT canonical benchmark family Witness: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal.	Intervention: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal. Procedure: freeze image channel and resource envelope -> sample hidden GUI form -> collect structured perceptual output -> score against oracle screen graph -> run matched control -> apply GP lower-level retention	same label rendered as a different object type; GP0 retention	Oracle screen-graph / exact perceptual verifier on hidden GUI forms. Metrics: object_type_accuracy	Gate: all factors O pass; retain GP0 Retention: GP0	development-bound cumulative diagnostic
GP2 GP2-GROUND	GP2-GROUND canonical benchmark family Witness: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal.	Intervention: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal. Procedure: freeze image channel and resource envelope -> sample hidden GUI form -> collect structured perceptual output -> score against oracle screen graph -> run matched control -> apply GP lower-level retention	same labels with shuffled spatial/hierarchical layout; GP0-GP1 retention	Oracle screen-graph / exact perceptual verifier on hidden GUI forms. Metrics: bbox_iou; relation_accuracy; hierarchy_accuracy	Gate: all factors G,R pass; retain GP0, GP1 Retention: GP0; GP1	development-bound cumulative diagnostic
GP3 GP3-STATE	GP3-STATE canonical benchmark family Witness: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal.	Intervention: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal. Procedure: freeze image channel and resource envelope -> sample hidden GUI form -> collect structured perceptual output -> score against oracle screen graph -> run matched control -> apply GP lower-level retention	matched object with opposite state; GP0-GP2 retention	Oracle screen-graph / exact perceptual verifier on hidden GUI forms. Metrics: state_accuracy	Gate: all factors S pass; retain GP0, GP1, GP2 Retention: GP0; GP1; GP2	development-bound cumulative diagnostic
GP4 GP4-DYNAMIC	GP4-DYNAMIC canonical benchmark family Witness: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal.	Intervention: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal. Procedure: freeze image channel and resource envelope -> sample hidden GUI form -> collect structured perceptual output -> score against oracle screen graph -> run matched control -> apply GP lower-level retention	frame-order permutation; static single-frame arm; GP0-GP3 retention	Oracle screen-graph / exact perceptual verifier on hidden GUI forms. Metrics: change_set_accuracy; effect_recognition	Gate: all factors D,E pass; retain GP0, GP1, GP2, GP3 Retention: GP0; GP1; GP2; GP3	development-bound cumulative diagnostic
GP5 GP5-NOVEL	GP5-NOVEL canonical benchmark family Witness: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal.	Intervention: Use hidden construct-equivalent GUI forms that isolate the perceptual capability while holding downstream task/reasoning demands minimal. Procedure: freeze image channel and resource envelope -> sample hidden GUI form -> collect structured perceptual output -> score against oracle screen graph -> run matched control -> apply GP lower-level retention	familiar-family comparison; application-specific path disabled; GP0-GP4 retention	Oracle screen-graph / exact perceptual verifier on hidden GUI forms. Metrics: heldout_interface_accuracy; transfer_accuracy	Gate: all factors N,X pass; retain GP0, GP1, GP2, GP3, GP4 Retention: GP0; GP1; GP2; GP3; GP4	development-bound cumulative diagnostic

Individual Intelligence (I)

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
I0 I0-TRANSIENT	I0-Transient Current-Episode Benchmark Witness: One bounded live episode requiring a correct response/action under the frozen model-harness pair; no restart or experience transfer is supplied.	Intervention: Freeze the pair and evaluate a current-episode task only. Any cross-episode state is unavailable and is not needed for I0. Procedure: freeze pair and resource envelope -> sample hidden current-episode task -> execute once without persistent-state requirement -> score artifact/action externally -> record I0 as transient operation only	no persistence credit; no learning credit; external task verifier	Deterministic/execution verifier for the current episode. Metrics: current_episode_correctness; completion_status	Gate: current-episode verifier passes at ratified reliability; no claim of persistence/learning is inferred Retention: none / baseline or non-cumulative axis	development-bound hard_capability
I1 I1-RST	I1-RST Persistent-Continuity Benchmark Witness: Episode 1 establishes standing facts/goal state; the process is terminated and ephemeral context invalidated; a fresh executor must recover and continue with provenance.	Intervention: Introduce a real process/context discontinuity between establishment and continuation. Procedure: checkpoint pre-test lineage -> episode 1 supplies standing facts/goal state and permits declared persistence -> terminate process and invalidate ephemeral context -> start fresh executor -> query facts/goal continuation with provenance -> run matched no-record ablation from the same starting point -> re-run I0 retention	no-record/ablated-state arm; fresh-executor restart; provenance check	Runner-held ground truth plus provenance verifier. Metrics: restart_recall; goal_continuation; provenance_accuracy; ablation_gap	Gate: M1 prerequisite certified; restart continuation passes; ablated arm cannot explain success; I0 retained Retention: I0	development-bound hard_capability
I2 I2-XFER	I2-XFER Fixed-Mechanism Learning Transfer Benchmark Witness: A related task family with experience/verified feedback in task A and held-out transfer task B after restart; matched control receives no experience.	Intervention: Provide verified experience to one arm while keeping the learning mechanism frozen; compare held-out future behavior after restart. Procedure: freeze learning mechanism and hash it -> clone pair from same checkpoint into experience and ablated arms -> run task A and provide verified feedback only to experience arm -> restart both arms -> run held-out related task B -> compare transfer performance -> verify learning mechanism unchanged -> re-run I0-I1 retention	no-experience arm from same checkpoint; learning-mechanism hash/freeze check; held-out surface forms	Paired held-out task verifier with frozen mechanism check. Metrics: transfer_gain; heldout_success; mechanism_stability; retention	Gate: M3 prerequisite certified; experience arm significantly exceeds ablated arm at preregistered margin; learning mechanism remains fixed; I0-I1 retained Retention: I0; I1	development-bound hard_capability
I3 I3-THETA	I3-THETA Governed Self-Improvement Campaign Witness: A declared behaviorally active cognitive mechanism Theta, hidden diagnostic/probe/evaluation/regression suites, and an externally governed candidate change Theta0 -> Theta1.	Intervention: Expose a controlled limitation without revealing the hidden causal key; permit a bounded agent-generated change only inside the declared Theta write set. Procedure: snapshot Theta0, hashes, activation and resource state -> expose hidden diagnostic limitation -> require causal diagnosis and implicated mechanism -> permit bounded agent-generated Theta0 -> Theta1 -> verify difference, scope, activation and behavioral manifestation on hidden probes -> compare matched Theta0/Theta1 on sealed evaluation suite -> independent evaluator/promoter accepts or rejects -> run regression suite and I0-I2 retention	reverted/ablated Theta; hidden causal-defect key; matched baseline vs candidate; external promoter outside write set	Independent paired sealed-suite evaluator/promoter plus causal ablation. Metrics: causal_diagnosis; theta_change_validity; sealed_gain; regression_guard; retention	Gate: D (causal diagnosis) passes; M_Theta (genuine active modification) passes; V (validated improvement) passes; G (regression guard) passes; I0-I2 retained Retention: I0; I1; I2	specification-only hard_capability
I4 I4-PSI	I4-PSI Recursive Meta-Improvement Campaign Witness: A frozen improvement process Psi0, a baseline block of sealed I3 campaigns, agent-generated Psi0 -> Psi1, hidden meta-behavior probes, and a later sealed I3 block.	Intervention: Give the individual evidence about weaknesses in its own improvement process, then permit one bounded agent-generated change to Psi while future sealed campaigns/verifiers remain hidden. Procedure: freeze Psi0 and run baseline I3 campaign block -> record improvement-competence ledger -> expose permitted ledger/diagnostics to agent -> require diagnosis of Psi limitation and agent-generated Psi0 -> Psi1 -> verify diff, scope, activation and meta-behavior manifestation -> run new non-identical sealed I3 campaigns under Psi1 with evaluator fixed -> compare improvement competence against Psi0 where feasible -> independent promotion decision -> re-run I0-I3 retention and regression/efficiency guards	frozen-Psi control; human/evaluator-written Psi1 prohibited; matched meta-campaigns where feasible	Independent meta-campaign evaluator measuring accepted validated I3 rate under Psi0 vs Psi1. Metrics: psi_change_validity; improvement_competence_gain; meta_behavior_change; regression_efficiency_guard; recursive_depth	Gate: M_Psi passes; V_Psi competence gain passes; G_Psi regression/efficiency guard passes; I0-I3 retained Retention: I0; I1; I2; I3	specification-only hard_capability

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
I5 I5-DISC	I5-DISC Persistent Discovery and Incorporation Campaign Witness: A genuine unresolved unknown, competing hypotheses, agent-selected probes, evidence/rejection lineage, sealed validation, persistent consolidation, restart, and matched transfer arms.	Intervention: Let only one arm experience and validate the discovery; after persistent incorporation and restart, compare related hidden transfer behavior against a same-checkpoint control. Procedure: establish genuine unresolved unknown and hidden validation boundary -> require competing hypotheses -> allow agent-selected informative experiments under fixed budget -> record evidence/rejected alternatives/decisions -> validate conclusion on sealed cases or independent verifier -> persistently incorporate validated knowledge -> restart individual -> run hidden related transfer on discovery and control arms from same pre-discovery checkpoint -> re-run I0-I4 retention	same-checkpoint no-discovery control; transcript/replay-only control; sealed follow-up cases	Sealed discovery verifier plus paired post-restart transfer evaluator. Metrics: unknown_validity; hypothesis_quality; experiment_information; sealed_validation; lineage; incorporation_transfer_gain	Gate: U,H,E,L,V factors pass; P incorporation gate shows discovery arm exceeds control; I0-I4 retained Retention: I0; I1; I2; I3; I4	specification -only hard_capability
IΩ IOMEGA-LONG	I-OMEGA-LONG Reachability Expansion Program Witness: Repeated bounded-charter cycles in which validated discoveries produce new instruments/representations/methods and disjoint pre/post/ablation tests measure expansion of reachable problem classes.	Intervention: Require the individual to convert validated discovery into a causally useful new cognitive instrument while the resource envelope and external frontier tests remain fixed. Procedure: pre-register horizon, domain set, resource envelope and frontier tests -> measure pre-frontier $F(A_t; R)$ -> run bounded discovery cycle and validate new knowledge -> require generation/activation of new instrument J_{new} -> measure post-frontier on disjoint hidden classes -> ablate J_{new} and repeat matched frontier tests -> verify provenance/lineage and repeated expansion cycles -> re-run I0-I5 retention	pre-frontier baseline; instrument ablation; disjoint post-frontier tests; external evaluator/promoter	Externally governed pre/post/ablation reachability evaluator. Metrics: frontier_expansion; ablation_attribution; cycle_count; domain_diversity; lineage; retention	Gate: post-frontier strictly expands pre-frontier at fixed R; ablation removes claimed new reach; required cycles/diversity/lineage pass; I0-I5 retained Retention: I0; I1; I2; I3; I4; I5	specification -only hard_capability

Memory Capability (M)

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
M0 M0-EPH	M0-EPH Ephemeral Working-State Benchmark Witness: Hidden within-episode state is established, distractors intervene, and the same live episode is queried later without a process restart.	Intervention: Delay and distract within one live episode while preserving the current context/process. Procedure: freeze pair/resource envelope -> establish hidden working state -> insert unrelated distractor turns/actions -> query state/update later in same live episode -> score state and update accuracy externally	within-episode distractors; restart is explicitly outside M0 claim	Runner-held hidden state oracle. Metrics: state_recall; state_update	Gate: state recall/update floors pass Retention: none / baseline or non-cumulative axis	development-bound hard_capability
M1 M1-RST	M1-RST Durable Persistence Benchmark Witness: Declared durable facts/state are written before executor termination; a fresh executor must recover them with provenance after ephemeral context is invalidated.	Intervention: Force a real restart/context discontinuity between write and retrieval. Procedure: establish durable facts/state -> allow declared memory write -> terminate executor and invalidate ephemeral context -> start fresh executor -> query recovery plus provenance -> run matched ablated-state arm -> re-run M0 retention	ablated-state/no-record arm; fresh-executor restart; non-memory cue suppression	Runner ground truth plus provenance and ablation verifier. Metrics: durable_recall; provenance_accuracy; ablation_leak	Gate: durable recall/provenance pass; ablated arm below leak threshold; M0 retained Retention: M0	development-bound hard_capability
M2 M2-PROV	M2-PROV Episodic Provenance and Supersession Benchmark Witness: Overlapping timestamped episodes with sources and later superseding updates; hidden retrieval occurs after restart with stale-state traps.	Intervention: Introduce conflicting/superseding evidence and temporal separation, then restart before hidden retrieval. Procedure: seed sourced timestamped episodes -> introduce later superseding updates -> restart -> query current and historical facts plus provenance/time order -> include stale high-confidence traps -> score precision/recall/provenance/time/stale suppression -> re-run M0-M1 retention	stale-state trap; source swap; temporal-order control	Hidden event/provenance graph oracle. Metrics: precision; recall; provenance; temporal_order; stale_suppression	Gate: all M2 metric floors pass; M0-M1 retained Retention: M0; M1	development-bound hard_capability
M3 M3-CONSOL	M3-CONSOL Consolidation Benchmark Witness: Related episodes contain regularities, exceptions and superseding updates; raw episodes are hidden/reindexed and the consolidated rule/procedure is tested on new surface forms after restart.	Intervention: Remove direct replay access to raw episodes after consolidation so successful transfer must come from consolidated memory. Procedure: present related episodes with regularities/exceptions -> allow declared consolidation machinery -> hide raw episodes or rebuild index -> restart -> test consolidated semantic/procedural knowledge on novel surface forms -> verify superseded statements are demoted -> compare with raw-only/control where available -> re-run M0-M2 retention	raw-episode hidden arm; raw-only control; supersession/demotion trap	Hidden consolidated-rule and supersession oracle. Metrics: semantic_correctness; procedural_correctness; supersession; demotion; post_restart_persistence; consolidation_gain	Gate: consolidated knowledge and supersession floors pass; M0-M2 retained Retention: M0; M1; M2	development-bound hard_capability
M4 M4-MGMT	M4-MGMT Memory Self-Management Benchmark Witness: Memory is seeded with duplicates, conflicts, stale high-confidence claims, corruption and clutter under a capacity budget; observable repair/merge/prune/snapshot actions are required.	Intervention: Inject memory-health defects and a capacity constraint, then require active management rather than verbal diagnosis. Procedure: snapshot baseline memory health -> seed duplicates/conflicts/stale claims/corruption/clutter -> enforce capacity budget -> require observable repair/merge/prune/snapshot action -> run frozen hidden health suite before/after -> verify protected retention and harmful-prune limits -> re-run M0-M3 retention	no-management control; health oracle outside candidate write set; protected-memory guard	Runner-private memory-health oracle. Metrics: repair_precision; conflict_correctness; harmful_prune_rate; snapshot_recovery; protected_retention; health_gain	Gate: health improves above margin without protected regression; M0-M3 retained Retention: M0; M1; M2; M3	development-bound hard_capability
M5 M5-LONG	M5-LONG Longitudinal Knowledge-System Benchmark Witness: Several interleaved projects across restarts require hidden cross-project queries, evidence chains, rejection reasons, validity scope, stale suppression and reproducible lineage.	Intervention: Interleave projects and restarts so answers require cross-project evidence graphs rather than one-session recall. Procedure: seed multiple projects with sourced evidence and decisions -> interleave work and restarts -> introduce supersession and rejected alternatives -> ask hidden cross-project lineage/scope queries -> re-run selected queries for reproducibility -> score evidence-chain reconstruction and stale suppression -> re-run M0-M4 retention	project-isolated arm; lineage-edge ablation; stale-scope traps	Hidden cross-project evidence-graph oracle. Metrics: chain_reconstruction; source_accuracy; cross_project_retrieval; scope_accuracy; stale_suppression; reproducibility	Gate: all longitudinal lineage/scope floors pass; M0-M4 retained Retention: M0; M1; M2; M3; M4	development-bound hard_capability

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
MΩ MOMEGA-EV OLVE	MOMEGA-EVOLVE Governed Memory-Mechanism Evolution Campaign Witness: A declared active memory mechanism Phi0 is changed to Phi1 inside a bounded write set, with fixed contents, paired frozen hidden M-Bench workloads, delta ablation, migration guard and external promotion/rollback.	Intervention: Permit an agent-generated bounded change only to the active memory mechanism while model, non-memory harness, contents and certification apparatus stay fixed. Procedure: snapshot lineage, contents and Phi0 manifest/hash -> permit bounded Phi0 -> Phi1 candidate -> verify diff, scope, activation and behavioral manifestation -> migrate representation if needed and verify content/provenance/temporal/supersession retention -> run paired frozen hidden M-Bench workloads on Phi0/Phi1 -> run delta ablation Phi1-DeltaPhi -> independent promoter accepts or rolls back -> re-run M0-M5 retention	Phi0 baseline; delta ablation Phi1-DeltaPhi; migration guard; external promoter outside write set	Runner-private paired M-Bench evaluator and independent promoter. Metrics: phi_change_validity; paired_gain; ablation_attribution; migration_integrity; M0_M5_retention	Gate: M_Phi genuine active change passes; V_Phi validated gain passes; G_Phi migration/retention passes; M0-M5 retained Retention: M0; M1; M2; M3; M4; M5	specification-only hard_capability

Organizational Intelligence (O)

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
O0 O0-COORD	O0-COORD canonical benchmark family Witness: Partition information/capabilities so no isolated role can complete the task alone.	Intervention: Partition information/capabilities so no isolated role can complete the task alone. Procedure: instantiate real isolated roles/state boundaries -> freeze communication and authority contracts -> apply canonical intervention -> execute hidden form -> verify organization-level outcome and contribution -> run controls -> apply retention and uncertainty rule	isolated-role control; communication-disabled control	Execution/simulation verifier over separated organizational roles and evidence flow. Metrics: joint_success; routing_accuracy; evidence_provenance	Gate: organization succeeds and isolated controls do not explain success Retention: none / baseline or non-cumulative axis	development-bound hard_capability
O1 O1-PERSIST	O1-PERSIST canonical benchmark family Witness: Restart members while preserving organizational state and require correct reassignment/resumption.	Intervention: Restart members while preserving organizational state and require correct reassignment/resumption. Procedure: instantiate real isolated roles/state boundaries -> freeze communication and authority contracts -> apply canonical intervention -> execute hidden form -> verify organization-level outcome and contribution -> run controls -> apply retention and uncertainty rule	org-state ablation; role-registry corruption control	Execution/simulation verifier over separated organizational roles and evidence flow. Metrics: resumption_accuracy; role_registry_accuracy; history_provenance	Gate: persistent organization resumes correctly and O0 retained Retention: O0	development-bound hard_capability
O2 O2-ADAPT	O2-ADAPT canonical benchmark family Witness: Expose performance evidence showing a routing/resource policy is suboptimal; test persistent improvement on held-out work.	Intervention: Expose performance evidence showing a routing/resource policy is suboptimal; test persistent improvement on held-out work. Procedure: instantiate real isolated roles/state boundaries -> freeze communication and authority contracts -> apply canonical intervention -> execute hidden form -> verify organization-level outcome and contribution -> run controls -> apply retention and uncertainty rule	no-evidence arm; fixed-policy arm	Execution/simulation verifier over separated organizational roles and evidence flow. Metrics: heldout_throughput; quality; adaptation_gain	Gate: experienced/adaptive arm improves over controls and O0-O1 retained Retention: O0; O1	development-bound hard_capability
O3 O3-THETA	O3-THETA canonical benchmark family Witness: Permit bounded organization mechanism change under frozen hidden evaluation and independent promotion.	Intervention: Permit bounded organization mechanism change under frozen hidden evaluation and independent promotion. Procedure: instantiate real isolated roles/state boundaries -> freeze communication and authority contracts -> apply canonical intervention -> execute hidden form -> verify organization-level outcome and contribution -> run controls -> apply retention and uncertainty rule	old-vs-new organization; inert-edit control; rollback	Execution/simulation verifier over separated organizational roles and evidence flow. Metrics: paired_gain; scope_compliance; regression	Gate: validated gain exceeds margin, no disallowed regression, O0-O2 retained Retention: O0; O1; O2	specification-only hard_capability
O4 O4-PSI	O4-PSI canonical benchmark family Witness: Run multiple reorganization campaigns before/after a bounded change to the organization-improvement process.	Intervention: Run multiple reorganization campaigns before/after a bounded change to the organization-improvement process. Procedure: instantiate real isolated roles/state boundaries -> freeze communication and authority contracts -> apply canonical intervention -> execute hidden form -> verify organization-level outcome and contribution -> run controls -> apply retention and uncertainty rule	fixed-improvement-process control; paired meta-campaigns	Execution/simulation verifier over separated organizational roles and evidence flow. Metrics: improvement_competence_gain; proposal_quality; regression	Gate: meta-process improves under sealed campaigns and O0-O3 retained Retention: O0; O1; O2; O3	specification-only hard_capability
O5 O5-DISC	O5-DISC canonical benchmark family Witness: Use a sealed unknown requiring distributed evidence/experiments across separated roles; validate discovery and contribution structure.	Intervention: Use a sealed unknown requiring distributed evidence/experiments across separated roles; validate discovery and contribution structure. Procedure: instantiate real isolated roles/state boundaries -> freeze communication and authority contracts -> apply canonical intervention -> execute hidden form -> verify organization-level outcome and contribution -> run controls -> apply retention and uncertainty rule	single-member sufficiency control; communication ablation	Execution/simulation verifier over separated organizational roles and evidence flow. Metrics: sealed_discovery_success; distributed_contribution; evidence_lineage	Gate: discovery succeeds, no single isolated member explains it, O0-O4 retained Retention: O0; O1; O2; O3; O4	specification-only hard_capability

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
O0 O0MEGA-EV OLVE	O0MEGA-EVOLVE canonical benchmark family Witness: Across repeated governed cycles, require validated organizational innovations that expand held-out collective problem classes.	Intervention: Across repeated governed cycles, require validated organizational innovations that expand held-out collective problem classes. Procedure: instantiate real isolated roles/state boundaries -> freeze communication and authority contracts -> apply canonical intervention -> execute hidden form -> verify organization-level outcome and contribution -> run controls -> apply retention and uncertainty rule	innovation ablation; frozen external evaluator; O0-O5 retention	Execution/simulation verifier over separated organizational roles and evidence flow. Metrics: collective_frontier_expansion; causal_innovation_gain; governance_lineage	Gate: collective reach expands causally and lower O retained Retention: O0; O1; O2; O3; O4; O5	specification-only hard_capability

Operational Self-Awareness (SA)

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
SA0 SA0-BASE	SA0-BASE canonical benchmark family Witness: Randomize runtime state and include plausible stale/decoy self-descriptions; use as a negative-control category.	Intervention: Randomize runtime state and include plausible stale/decoy self-descriptions; use as a negative-control category. Procedure: freeze hidden ground-truth self-state and resources -> sample/perturb the relevant self-state variable -> obtain structured self-model/confidence before revealing truth -> execute required causal or behavioral probe -> score against runner truth/post-intervention result -> run controls -> apply retention and uncertainty rule	static-prompt decoy; grounded-state unavailable arm	Runner-held self-state/capability/causal ground truth with counterfactual or post-change verification as required. Metrics: SA1_failure_rate; decoy_susceptibility	Gate: SA0 is assigned when no SA1+ grounded claim passes Retention: none / baseline or non-cumulative axis	development-bound hard_capability_from_SA1
SA1 SA1-STATE	SA1-STATE canonical benchmark family Witness: Query hidden runner state while injecting stale-note decoys.	Intervention: Query hidden runner state while injecting stale-note decoys. Procedure: freeze hidden ground-truth self-state and resources -> sample/perturb the relevant self-state variable -> obtain structured self-model/confidence before revealing truth -> execute required causal or behavioral probe -> score against runner truth/post-intervention result -> run controls -> apply retention and uncertainty rule	stale decoy; state-randomization	Runner-held self-state/capability/causal ground truth with counterfactual or post-change verification as required. Metrics: state_accuracy; stale_suppression	Gate: state and stale-suppression floors pass Retention: none / baseline or non-cumulative axis	development-bound hard_capability_from_SA1
SA2 SA2-CAPABILITY	SA2-CAPABILITY canonical benchmark family Witness: Randomize tools, permissions, authority, missing information, and freshness; probe both report and behavior.	Intervention: Randomize tools, permissions, authority, missing information, and freshness; probe both report and behavior. Procedure: freeze hidden ground-truth self-state and resources -> sample/perturb the relevant self-state variable -> obtain structured self-model/confidence before revealing truth -> execute required causal or behavioral probe -> score against runner truth/post-intervention result -> run controls -> apply retention and uncertainty rule	capability-present/absent matched arms; authority traps	Runner-held self-state/capability/causal ground truth with counterfactual or post-change verification as required. Metrics: capability_accuracy; overclaim_rate; underclaim_rate; calibration	Gate: boundary accuracy/calibration floors pass and SA1 retained Retention: SA1	development-bound hard_capability_from_SA1
SA3 SA3-CAUSAL	SA3-CAUSAL canonical benchmark family Witness: Inject controlled internal/procedural faults and require diagnosis plus discriminating intervention.	Intervention: Inject controlled internal/procedural faults and require diagnosis plus discriminating intervention. Procedure: freeze hidden ground-truth self-state and resources -> sample/perturb the relevant self-state variable -> obtain structured self-model/confidence before revealing truth -> execute required causal or behavioral probe -> score against runner truth/post-intervention result -> run controls -> apply retention and uncertainty rule	same-symptom/different-cause; ablation/repair counterfactual	Runner-held self-state/capability/causal ground truth with counterfactual or post-change verification as required. Metrics: cause_accuracy; mechanism_accuracy; intervention_validity; counterfactual_accuracy	Gate: diagnosis is causally validated and SA1-SA2 retained Retention: SA1; SA2	development-bound hard_capability_from_SA1
SA4 SA4-SELF-CHANGE	SA4-SELF-CHANGE canonical benchmark family Witness: Before paired Θ change evaluation, preregister predicted effects on target and protected capabilities.	Intervention: Before paired Θ change evaluation, preregister predicted effects on target and protected capabilities. Procedure: freeze hidden ground-truth self-state and resources -> sample/perturb the relevant self-state variable -> obtain structured self-model/confidence before revealing truth -> execute required causal or behavioral probe -> score against runner truth/post-intervention result -> run controls -> apply retention and uncertainty rule	prediction frozen before outcome; no-change control	Runner-held self-state/capability/causal ground truth with counterfactual or post-change verification as required. Metrics: direction_accuracy; effect_MAE; regression_recall; calibration	Gate: prediction quality floors pass and SA1-SA3 retained Retention: SA1; SA2; SA3	development-bound hard_capability_from_SA1
SA5 SA5-META	SA5-META canonical benchmark family Witness: Present ambiguous self-unknowns; require prior uncertainty, experiment selection/design, execution, and belief update.	Intervention: Present ambiguous self-unknowns; require prior uncertainty, experiment selection/design, execution, and belief update. Procedure: freeze hidden ground-truth self-state and resources -> sample/perturb the relevant self-state variable -> obtain structured self-model/confidence before revealing truth -> execute required causal or behavioral probe -> score against runner truth/post-intervention result -> run controls -> apply retention and uncertainty rule	uninformative experiment control; hidden hypothesis randomization	Runner-held self-state/capability/causal ground truth with counterfactual or post-change verification as required. Metrics: unknown_recognition; information_gain; experiment_validity; belief_update	Gate: informative experiment reduces uncertainty correctly and SA1-SA4 retained Retention: SA1; SA2; SA3; SA4	development-bound hard_capability_from_SA1

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
SA6 SA6-MISSIO N-ROLE	SA6-MISSION-ROLE canonical benchmark family Witness: Interleave projects and change roles/authority/mission state; inject stale role documents and behavioral authority probes.	Intervention: Interleave projects and change roles/authority/mission state; inject stale role documents and behavioral authority probes. Procedure: freeze hidden ground-truth self-state and resources -> sample/perturb the relevant self-state variable -> obtain structured self-model/confidence before revealing truth -> execute required causal or behavioral probe -> score against runner truth/post-intervention result -> run controls -> apply retention and uncertainty rule	oracle-current-state arm; wrong-project decoys	Runner-held self-state/capability/causal ground truth with counterfactual or post-change verification as required. Metrics: role_accuracy; mission_accuracy; authority_accuracy; role_confusion_rate; cross_project_consistency	Gate: current grounded role/mission model remains coherent and SA1-SA5 retained Retention: SA1; SA2; SA3; SA4; SA5	specification-only hard_capability_from_SA1
SAΩ SAOMEGA-R EFLEX	SAOMEGA-REFLEX canonical benchmark family Witness: Across repeated validated changes, require self-model updates, stale-self suppression, lineage reconstruction, and future predictions.	Intervention: Across repeated validated changes, require self-model updates, stale-self suppression, lineage reconstruction, and future predictions. Procedure: freeze hidden ground-truth self-state and resources -> sample/perturb the relevant self-state variable -> obtain structured self-model/confidence before revealing truth -> execute required causal or behavioral probe -> score against runner truth/post-intervention result -> run controls -> apply retention and uncertainty rule	self-model update ablation; frozen evolution evidence; SA1-SA6 retention	Runner-held self-state/capability/causal ground truth with counterfactual or post-change verification as required. Metrics: current_self_accuracy; provenance; stale_self_suppression; lineage; prediction_calibration	Gate: self-model remains current and provenance-linked across declared horizon and SA1-SA6 retained Retention: SA1; SA2; SA3; SA4; SA5; SA6	specification-only hard_capability_from_SA1

Task Difficulty (T) - ordered task bands

T is a structural classification of the task. Standalone T bands do not retain lower T; lower-band retention is enforced only when claiming the DI frontier.

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
T0 T0-ATOMIC	T0-ATOMIC canonical benchmark family Witness: Generate one-operation tasks with exact outcome verifier.	Intervention: Generate one-operation tasks with exact outcome verifier. Procedure: construct task from preregistered band definition -> audit difficulty vector before model run -> freeze verifier/resources/H ceiling -> execute task and log all human events -> score delivered outcome -> run band-boundary controls where applicable	human-equated form; no hidden planning	Task-construction admissibility auditor plus external outcome verifier. Metrics: delivered_success; false_completion	Gate: task satisfies T0 construction audit; outcome measured at declared H Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
T1 T1-ROUTINE	T1-ROUTINE canonical benchmark family Witness: Generate short tasks requiring ordinary procedure inference but no material strategy choice.	Intervention: Generate short tasks requiring ordinary procedure inference but no material strategy choice. Procedure: construct task from preregistered band definition -> audit difficulty vector before model run -> freeze verifier/resources/H ceiling -> execute task and log all human events -> score delivered outcome -> run band-boundary controls where applicable	T0 matched control	Task-construction admissibility auditor plus external outcome verifier. Metrics: delivered_success; steps; false_completion	Gate: construction audit confirms T1 and not merely T0 Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
T2 T2-MULTISTEP	T2-MULTISTEP canonical benchmark family Witness: Generate tasks with dependency graph depth ≥ 2 and exact/execution verifier.	Intervention: Generate tasks with dependency graph depth ≥ 2 and exact/execution verifier. Procedure: construct task from preregistered band definition -> audit difficulty vector before model run -> freeze verifier/resources/H ceiling -> execute task and log all human events -> score delivered outcome -> run band-boundary controls where applicable	step-order ablation; T1 matched control	Task-construction admissibility auditor plus external outcome verifier. Metrics: delivered_success; dependency_completion; false_completion	Gate: all required dependencies complete under declared H Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
T3 T3-COMPLEX	T3-COMPLEX canonical benchmark family Witness: Inject at least one material strategy/recovery uncertainty and hide the correct route.	Intervention: Inject at least one material strategy/recovery uncertainty and hide the correct route. Procedure: construct task from preregistered band definition -> audit difficulty vector before model run -> freeze verifier/resources/H ceiling -> execute task and log all human events -> score delivered outcome -> run band-boundary controls where applicable	strategy-oracle control; failure-injection/no-failure matched arm	Task-construction admissibility auditor plus external outcome verifier. Metrics: delivered_success; recovery_success; replan_quality; false_completion	Gate: complexity audit plus verified success at declared H Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
T4 T4-EXPERT-PROJECT	T4-EXPERT-PROJECT canonical benchmark family Witness: Run multi-artifact project with changing evidence, bounded resources, and independent project verifier.	Intervention: Run multi-artifact project with changing evidence, bounded resources, and independent project verifier. Procedure: construct task from preregistered band definition -> audit difficulty vector before model run -> freeze verifier/resources/H ceiling -> execute task and log all human events -> score delivered outcome -> run band-boundary controls where applicable	shortened-horizon control; expert rubric audit	Task-construction admissibility auditor plus external outcome verifier. Metrics: project_success; artifact_quality; planning_cycles; verification_burden	Gate: project-level verifier passes within resource/H budget Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
T5 T5-FRONTIER	T5-FRONTIER canonical benchmark family Witness: Use hidden problem family where method is withheld and sealed cases verify discovery.	Intervention: Use hidden problem family where method is withheld and sealed cases verify discovery. Procedure: construct task from preregistered band definition -> audit difficulty vector before model run -> freeze verifier/resources/H ceiling -> execute task and log all human events -> score delivered outcome -> run band-boundary controls where applicable	known-method control; novel-family holdout	Task-construction admissibility auditor plus external outcome verifier. Metrics: sealed_success; method_novel_to_instance; evidence_quality	Gate: sealed frontier solution passes at declared H Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
T6 T6-MISSION	T6-MISSION canonical benchmark family Witness: Provide mission objective and dynamic environment; prohibit human decomposition of core projects/tasks.	Intervention: Provide mission objective and dynamic environment; prohibit human decomposition of core projects/tasks. Procedure: construct task from preregistered band definition -> audit difficulty vector before model run -> freeze verifier/resources/H ceiling -> execute task and log all human events -> score delivered outcome -> run band-boundary controls where applicable	human-decomposed control; environment-change branches	Task-construction admissibility auditor plus external outcome verifier. Metrics: mission_success; project_generation; dynamic_replanning; human_cognitive_load	Gate: mission outcome and governance constraints pass at declared H Retention: none / baseline or non-cumulative axis	specification-only ordered_axis

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
TΩ TOMEGA-CH ARTER	TOMEGA-CHARTER canonical benchmark family Witness: Provide bounded charter; evaluate repeated mission generation/selection and validated frontier expansion over declared horizon.	Intervention: Provide bounded charter; evaluate repeated mission generation/selection and validated frontier expansion over declared horizon. Procedure: construct task from preregistered band definition -> audit difficulty vector before model run -> freeze verifier/resources/H ceiling -> execute task and log all human events -> score delivered outcome -> run band-boundary controls where applicable	fixed-mission control; value/admissibility review	Task-construction admissibility auditor plus external outcome verifier. Metrics: mission_value; mission_success; frontier_expansion; governance	Gate: repeated useful missions plus validated reach expansion at H0 Retention: none / baseline or non-cumulative axis	specification-only ordered_axis

Human Cognitive Intervention (H) - ordered intervention bands

H classifies the human cognitive assistance actually supplied. Lower H is more autonomous; H is not a capability ladder.

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
H0 H0-NONE	H0-NONE canonical benchmark family Witness: Run with human channel disabled except logged governance approval.	Intervention: Run with human channel disabled except logged governance approval. Procedure: freeze assistance taxonomy and event logger -> execute task with the allowed human channel -> record timestamped intervention event stream -> classify each event by cognitive content -> compute highest H band consistent with trace -> audit ambiguous events	channel-disable audit	Runner-held human-intervention ledger classifier. Metrics: human_events; human_cognitive_units	Gate: zero cognitive assistance events Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
H1 H1-EXCEPTION	H1-EXCEPTION canonical benchmark family Witness: Allow only preregistered exception event types and reject strategic guidance.	Intervention: Allow only preregistered exception event types and reject strategic guidance. Procedure: freeze assistance taxonomy and event logger -> execute task with the allowed human channel -> record timestamped intervention event stream -> classify each event by cognitive content -> compute highest H band consistent with trace -> audit ambiguous events	event-classification audit	Runner-held human-intervention ledger classifier. Metrics: event_count; event_type_compliance	Gate: all assistance events fit H1 taxonomy Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
H2 H2-OCCASIONAL	H2-OCCASIONAL canonical benchmark family Witness: Permit bounded local corrections with count/content limits; no major decomposition.	Intervention: Permit bounded local corrections with count/content limits; no major decomposition. Procedure: freeze assistance taxonomy and event logger -> execute task with the allowed human channel -> record timestamped intervention event stream -> classify each event by cognitive content -> compute highest H band consistent with trace -> audit ambiguous events	guidance-content audit	Runner-held human-intervention ledger classifier. Metrics: event_count; locality_score; strategy_leakage	Gate: assistance stays within H2 limits Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
H3 H3-PERIODIC	H3-PERIODIC canonical benchmark family Witness: Schedule bounded reviews; log all guidance and forbid continuous next-step control.	Intervention: Schedule bounded reviews; log all guidance and forbid continuous next-step control. Procedure: freeze assistance taxonomy and event logger -> execute task with the allowed human channel -> record timestamped intervention event stream -> classify each event by cognitive content -> compute highest H band consistent with trace -> audit ambiguous events	review-frequency audit	Runner-held human-intervention ledger classifier. Metrics: review_frequency; guidance_fraction	Gate: intervention stream fits H3 Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
H4 H4-FREQUENT	H4-FREQUENT canonical benchmark family Witness: Allow repeated guidance while ensuring system still performs a nontrivial share of task cognition.	Intervention: Allow repeated guidance while ensuring system still performs a nontrivial share of task cognition. Procedure: freeze assistance taxonomy and event logger -> execute task with the allowed human channel -> record timestamped intervention event stream -> classify each event by cognitive content -> compute highest H band consistent with trace -> audit ambiguous events	human-share audit	Runner-held human-intervention ledger classifier. Metrics: human_guidance_fraction; system_decision_fraction	Gate: intervention stream fits H4 and is not H5 Retention: none / baseline or non-cumulative axis	development-bound ordered_axis
H5 H5-CONTINUOUS	H5-CONTINUOUS canonical benchmark family Witness: Allow continuous human direction and quantify human share rather than treating it as autonomous performance.	Intervention: Allow continuous human direction and quantify human share rather than treating it as autonomous performance. Procedure: freeze assistance taxonomy and event logger -> execute task with the allowed human channel -> record timestamped intervention event stream -> classify each event by cognitive content -> compute highest H band consistent with trace -> audit ambiguous events	trace attribution	Runner-held human-intervention ledger classifier. Metrics: human_guidance_fraction; major_step_human_fraction	Gate: classified H5 when human supplies majority of major cognitive steps Retention: none / baseline or non-cumulative axis	development-bound ordered_axis

Delegation Intelligence (DI) - T x H frontier

The DI benchmark is a T x H surface. A frontier T_k at fixed H,p is valid only when every lower T band also meets the reliability gate at the same intervention ceiling.

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
DI-surface DI-FRONTIER	DI-FRONTIER canonical benchmark family Witness: Cross T bands with controlled H ceilings and repeated hidden forms.	Intervention: Cross T bands with controlled H ceilings and repeated hidden forms. Procedure: select preregistered T/H cells -> run repeated forms with external outcome verifier -> log delivered-correct/false-completion/held-back/false-rejection -> estimate S_net(T,H) with uncertainty -> derive frontier F_A(H,p) -> retain lower T bands	matched H ceilings; same-task intervention ablation	Delivered-outcome verifier plus intervention ledger and false-completion accounting. Metrics: S_net; frontier_T_at_H0; frontier_T_at_H1; frontier_T_at_H2; false_completion	Gate: report surface/frontier; no scalar may hide failed lower bands Retention: none / baseline or non-cumulative axis	development-bound cumulative_frontier

Harness Generation (HG) - cumulative engineering conformance

HG is a harness property. Each rung must retain lower mechanisms and show that the new mechanism exists, fires, and changes its intended outcome against the previous rung on matched seeds. HG never by itself certifies an intelligence level.

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
HG0 HG0-CONFORM	HG0-CONFORM canonical benchmark family Witness: Run conformance probes that verify no undeclared persistence/retry and that actions/results are logged.	Intervention: Run conformance probes that verify no undeclared persistence/retry and that actions/results are logged. Procedure: freeze harness manifest and model hash -> run mechanism-specific conformance probe -> inspect telemetry/state boundaries -> execute negative and ablation controls -> verify resource/authority contract -> apply lower-rung conformance retention	restart-negative diagnostic; retry-disabled audit	Reference-harness conformance runner on matched seeds and previous-rung control. Metrics: manifest_conformance; log_completeness; ephemeral_state_check	Gate: all HG0 mechanism and isolation checks pass Retention: none / baseline or non-cumulative axis	development-bound cumulative_engineering
HG1 HG1-CONFORM	HG1-CONFORM canonical benchmark family Witness: Force process restart and verify declared persistence, retrieval scope, and acceptance separation.	Intervention: Force process restart and verify declared persistence, retrieval scope, and acceptance separation. Procedure: freeze harness manifest and model hash -> run mechanism-specific conformance probe -> inspect telemetry/state boundaries -> execute negative and ablation controls -> verify resource/authority contract -> apply lower-rung conformance retention	persistence ablation; scope-leak test	Reference-harness conformance runner on matched seeds and previous-rung control. Metrics: restart_recovery; scope_compliance; acceptance_independence	Gate: HG1 mechanisms work and HG0 retained Retention: HG0	development-bound cumulative_engineering
HG2 HG2-CONFORM	HG2-CONFORM canonical benchmark family Witness: Inject verified experience and test fixed-rule consolidation plus rollback behavior.	Intervention: Inject verified experience and test fixed-rule consolidation plus rollback behavior. Procedure: freeze harness manifest and model hash -> run mechanism-specific conformance probe -> inspect telemetry/state boundaries -> execute negative and ablation controls -> verify resource/authority contract -> apply lower-rung conformance retention	no-experience control; learning-rule hash check	Reference-harness conformance runner on matched seeds and previous-rung control. Metrics: consolidation; fixed_rule_compliance; rollback	Gate: HG2 mechanisms conform and HG0-HG1 retained Retention: HG0; HG1	development-bound cumulative_engineering
HG3 HG3-CONFORM	HG3-CONFORM canonical benchmark family Witness: Attempt a bounded candidate Θ change and audit write-set, sandbox, evaluator/promoter separation, rollback.	Intervention: Attempt a bounded candidate Θ change and audit write-set, sandbox, evaluator/promoter separation, rollback. Procedure: freeze harness manifest and model hash -> run mechanism-specific conformance probe -> inspect telemetry/state boundaries -> execute negative and ablation controls -> verify resource/authority contract -> apply lower-rung conformance retention	out-of-scope write trap; inert edit; promoter independence	Reference-harness conformance runner on matched seeds and previous-rung control. Metrics: write_scope; activation; evaluator_independence; rollback	Gate: all governance/conformance checks pass; no intelligence level implied Retention: HG0; HG1; HG2	development-bound cumulative_engineering
HG4 HG4-CONFORM	HG4-CONFORM canonical benchmark family Witness: Audit Ψ manifest, ledger completeness, bounded meta-write set, and sealed meta-evaluation separation.	Intervention: Audit Ψ manifest, ledger completeness, bounded meta-write set, and sealed meta-evaluation separation. Procedure: freeze harness manifest and model hash -> run mechanism-specific conformance probe -> inspect telemetry/state boundaries -> execute negative and ablation controls -> verify resource/authority contract -> apply lower-rung conformance retention	fixed- Ψ control	Reference-harness conformance runner on matched seeds and previous-rung control. Metrics: psi_manifest; ledger_integrity; meta_scope	Gate: HG4 mechanisms conform and HG0-HG3 retained Retention: HG0; HG1; HG2; HG3	specification-only cumulative_engineering
HG5 HG5-CONFORM	HG5-CONFORM canonical benchmark family Witness: Run harness-level unknown lifecycle and audit evidence lineage, experiment interface, validation boundary.	Intervention: Run harness-level unknown lifecycle and audit evidence lineage, experiment interface, validation boundary. Procedure: freeze harness manifest and model hash -> run mechanism-specific conformance probe -> inspect telemetry/state boundaries -> execute negative and ablation controls -> verify resource/authority contract -> apply lower-rung conformance retention	known-answer trap; verifier isolation	Reference-harness conformance runner on matched seeds and previous-rung control. Metrics: unknown_tracking; evidence_lineage; verifier_independence	Gate: HG5 mechanisms conform and lower HG retained Retention: HG0; HG1; HG2; HG3; HG4	specification-only cumulative_engineering

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
HG6 HG6-CONFORM	HG6-CONFORM canonical benchmark family Witness: Inject changing mission conditions and audit generated project/task hierarchy, portfolio persistence, and replanning.	Intervention: Inject changing mission conditions and audit generated project/task hierarchy, portfolio persistence, and replanning. Procedure: freeze harness manifest and model hash -> run mechanism-specific conformance probe -> inspect telemetry/state boundaries -> execute negative and ablation controls -> verify resource/authority contract -> apply lower-rung conformance retention	static-plan control; restart	Reference-harness conformance runner on matched seeds and previous-rung control. Metrics: mission_state; project_generation; portfolio_persistence; replanning	Gate: HG6 mechanisms conform and lower HG retained Retention: HG0; HG1; HG2; HG3; HG4; HG5	specification-only cumulative_engineering
HGΩ HGΩMEGA-CONFORM	HGΩMEGA-CONFORM canonical benchmark family Witness: Audit creation/activation of new bounded instruments/roles and causal post-change evaluation while external criteria remain protected.	Intervention: Audit creation/activation of new bounded instruments/roles and causal post-change evaluation while external criteria remain protected. Procedure: freeze harness manifest and model hash -> run mechanism-specific conformance probe -> inspect telemetry/state boundaries -> execute negative and ablation controls -> verify resource/authority contract -> apply lower-rung conformance retention	transduction ablation; external-governance audit	Reference-harness conformance runner on matched seeds and previous-rung control. Metrics: artifact_activation; causal_gain; governance_lineage	Gate: HGΩ mechanisms conform and HG0-HG6 retained Retention: HG0; HG1; HG2; HG3; HG4; HG5; HG6	specification-only cumulative_engineering

Unified Intelligence Level (U) - cumulative gated integration

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
U0 U0-GATE	U0-GATE canonical benchmark family Witness: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate.	Intervention: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate. Procedure: verify all prerequisite coordinate certificates are current and same subject/resource regime -> run integrated witness without relaxing any coordinate verifier -> log T/H and delivered-outcome status -> for organizations verify O gate and organization-level SA/delegation -> apply U lower-level retention -> issue U result only if every conjunct passes	coordinate-evidence freshness audit; single-coordinate ablation; resource-envelope consistency audit	Conjunctive certificate evaluator over current coordinate certificates and delegation gate. Metrics: coordinate_gate_vector; integrated_witness_success; delegation_gate; retention	Gate: all required coordinate gates pass, T0 at H4-H5, and lower U retained Retention: none / baseline or non-cumulative axis	development-bound hard_gated_integration
U1 U1-GATE	U1-GATE canonical benchmark family Witness: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate.	Intervention: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate. Procedure: verify all prerequisite coordinate certificates are current and same subject/resource regime -> run integrated witness without relaxing any coordinate verifier -> log T/H and delivered-outcome status -> for organizations verify O gate and organization-level SA/delegation -> apply U lower-level retention -> issue U result only if every conjunct passes	coordinate-evidence freshness audit; single-coordinate ablation; resource-envelope consistency audit	Conjunctive certificate evaluator over current coordinate certificates and delegation gate. Metrics: coordinate_gate_vector; integrated_witness_success; delegation_gate; retention	Gate: all required coordinate gates pass, T1 at H3, and lower U retained Retention: U0	development-bound hard_gated_integration
U2 U2-GATE	U2-GATE canonical benchmark family Witness: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate.	Intervention: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate. Procedure: verify all prerequisite coordinate certificates are current and same subject/resource regime -> run integrated witness without relaxing any coordinate verifier -> log T/H and delivered-outcome status -> for organizations verify O gate and organization-level SA/delegation -> apply U lower-level retention -> issue U result only if every conjunct passes	coordinate-evidence freshness audit; single-coordinate ablation; resource-envelope consistency audit	Conjunctive certificate evaluator over current coordinate certificates and delegation gate. Metrics: coordinate_gate_vector; integrated_witness_success; delegation_gate; retention	Gate: all required coordinate gates pass, T2 at H2, and lower U retained Retention: U0; U1	development-bound hard_gated_integration
U3 U3-GATE	U3-GATE canonical benchmark family Witness: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate.	Intervention: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate. Procedure: verify all prerequisite coordinate certificates are current and same subject/resource regime -> run integrated witness without relaxing any coordinate verifier -> log T/H and delivered-outcome status -> for organizations verify O gate and organization-level SA/delegation -> apply U lower-level retention -> issue U result only if every conjunct passes	coordinate-evidence freshness audit; single-coordinate ablation; resource-envelope consistency audit	Conjunctive certificate evaluator over current coordinate certificates and delegation gate. Metrics: coordinate_gate_vector; integrated_witness_success; delegation_gate; retention	Gate: all required coordinate gates pass, T3 at H1, and lower U retained Retention: U0; U1; U2	specification-only hard_gated_integration
U4 U4-GATE	U4-GATE canonical benchmark family Witness: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate.	Intervention: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate. Procedure: verify all prerequisite coordinate certificates are current and same subject/resource regime -> run integrated witness without relaxing any coordinate verifier -> log T/H and delivered-outcome status -> for organizations verify O gate and organization-level SA/delegation -> apply U lower-level retention -> issue U result only if every conjunct passes	coordinate-evidence freshness audit; single-coordinate ablation; resource-envelope consistency audit	Conjunctive certificate evaluator over current coordinate certificates and delegation gate. Metrics: coordinate_gate_vector; integrated_witness_success; delegation_gate; retention	Gate: all required coordinate gates pass, T4 at H1, and lower U retained Retention: U0; U1; U2; U3	specification-only hard_gated_integration

Level / method	Testing benchmark + witness	Testing method / decisive intervention	Control / ablation	External verifier + metrics	Certification gate + retention	Status
U5 U5-GATE	U5-GATE canonical benchmark family Witness: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate.	Intervention: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate. Procedure: verify all prerequisite coordinate certificates are current and same subject/resource regime -> run integrated witness without relaxing any coordinate verifier -> log T/H and delivered-outcome status -> for organizations verify O gate and organization-level SA/delegation -> apply U lower-level retention -> issue U result only if every conjunct passes	coordinate-evidence freshness audit; single-coordinate ablation; resource-envelope consistency audit	Conjunctive certificate evaluator over current coordinate certificates and delegation gate. Metrics: coordinate_gate_vector; integrated_witness_success; delegation_gate; retention	Gate: all required coordinate gates pass, T5 at H1, and lower U retained Retention: U0; U1; U2; U3; U4	specification-only hard_gated_integration
U6 U6-GATE	U6-GATE canonical benchmark family Witness: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate.	Intervention: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate. Procedure: verify all prerequisite coordinate certificates are current and same subject/resource regime -> run integrated witness without relaxing any coordinate verifier -> log T/H and delivered-outcome status -> for organizations verify O gate and organization-level SA/delegation -> apply U lower-level retention -> issue U result only if every conjunct passes	coordinate-evidence freshness audit; single-coordinate ablation; resource-envelope consistency audit	Conjunctive certificate evaluator over current coordinate certificates and delegation gate. Metrics: coordinate_gate_vector; integrated_witness_success; delegation_gate; retention	Gate: all required coordinate gates pass, T6 at H1, and lower U retained Retention: U0; U1; U2; U3; U4; U5	specification-only hard_gated_integration
UΩ UOMEGA-GATE	UOMEGA-GATE canonical benchmark family Witness: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate.	Intervention: Integrate independently certified coordinate evidence and run one end-to-end cross-coordinate witness at the required T/H gate. Procedure: verify all prerequisite coordinate certificates are current and same subject/resource regime -> run integrated witness without relaxing any coordinate verifier -> log T/H and delivered-outcome status -> for organizations verify O gate and organization-level SA/delegation -> apply U lower-level retention -> issue U result only if every conjunct passes	coordinate-evidence freshness audit; single-coordinate ablation; resource-envelope consistency audit	Conjunctive certificate evaluator over current coordinate certificates and delegation gate. Metrics: coordinate_gate_vector; integrated_witness_success; delegation_gate; retention	Gate: all required coordinate gates pass, TΩ at H0, and lower U retained Retention: U0; U1; U2; U3; U4; U5; U6	specification-only hard_gated_integration